

COL380

Introduction to
Parallel & Distributed Programming
2-0-2

Subodh Kumar

Agenda

- OS basic
- A review of Architecture

- Process & thread
 - ➔ Scheduling, Context-switching and concurrency
 - ➔ User space, Kernel space
 - ➔ System calls
- Address space & Name space
- Virtual Memory
 - ➔ Caches
- Synchronization

- ◆ Shell (bash)
- ◆ time
- ◆ PMU (perf)

Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)



Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)



1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

If TLB miss:

Walk Page Table, Add to TLB

Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

If TLB miss:

Walk Page Table, Add to TLB

Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

If TLB miss:

Walk Page Table, Add to TLB

If Page-table not resident in Memory:

Page Fault ⇒ Interrupt ⇒ Switch to Kernel

Get rescheduled after Page table locked in

Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

If TLB miss:
Walk Page Table, Add to TLB

If Page-table not resident in Memory:
Page Fault ⇒ Interrupt ⇒ Switch to Kernel
Get rescheduled after Page table locked in



Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

If TLB miss:
Walk Page Table, Add to TLB

If Page-table not resident in Memory:
Page Fault ⇒ Interrupt ⇒ Switch to Kernel
Get rescheduled after Page table locked in



Execution Steps

edit file.c:

```
#include ...  
void func(int a) {  
    sleep(a);  
}
```

```
int main(int args, char *argv[]) {  
    func(atoi(argv[0]));  
    ...  
    return v;  
}
```

gcc file.c -o exec

time ./exec

echo \$?

movl %eax, -4(%rbp)

1. Map Virtual addr → Physical
2. Fetch: L1, L2 .., Memory

(Multi-level) TLB-cache (CAM)
Get Physical Addr, Check protection

If TLB miss:
Walk Page Table, Add to TLB

If **page** not resident in Memory:
Page Fault ⇒ Interrupt ⇒ Switch to Kernel
Get rescheduled after Page table locked in



It's the Memory Stupid!

- Accessing variables is ^{much} 'slower' than computing
- Memory sizes are often sufficient
 - ➔ Particularly, when distributed among many nodes
 - ▶ Out-of-core computation needed sometimes (we won't focus)
- Caches are small
 - ➔ Not very helpful if traversing large swaths of memory
 - ➔ But, mind the line
 - ▶ And, prefetching is prevalent
 - ➔ Beware of sharing across caches, particularly false-sharing

Cache Example

```
for(int i=0; i<n; i++) {  
    a[i] = b[i] + c[i];  
}  
for(int i=0; i<n; i++) {  
    d[i] = e[i] + f[i];  
}
```

vs

```
for(int i=0; i<n; i++) {  
    a[i] = b[i] + c[i];  
    d[i] = e[i] + f[i];  
}
```

```
for(int i=0; i<N; i++) {  
    d[i] = x;  
}
```

```
for(int i=0; i<N*STEP; i++) {  
    d[(i*STEP)] = x;  
}
```

- Processes
- Address spaces
- Virtual memory
- Caches