

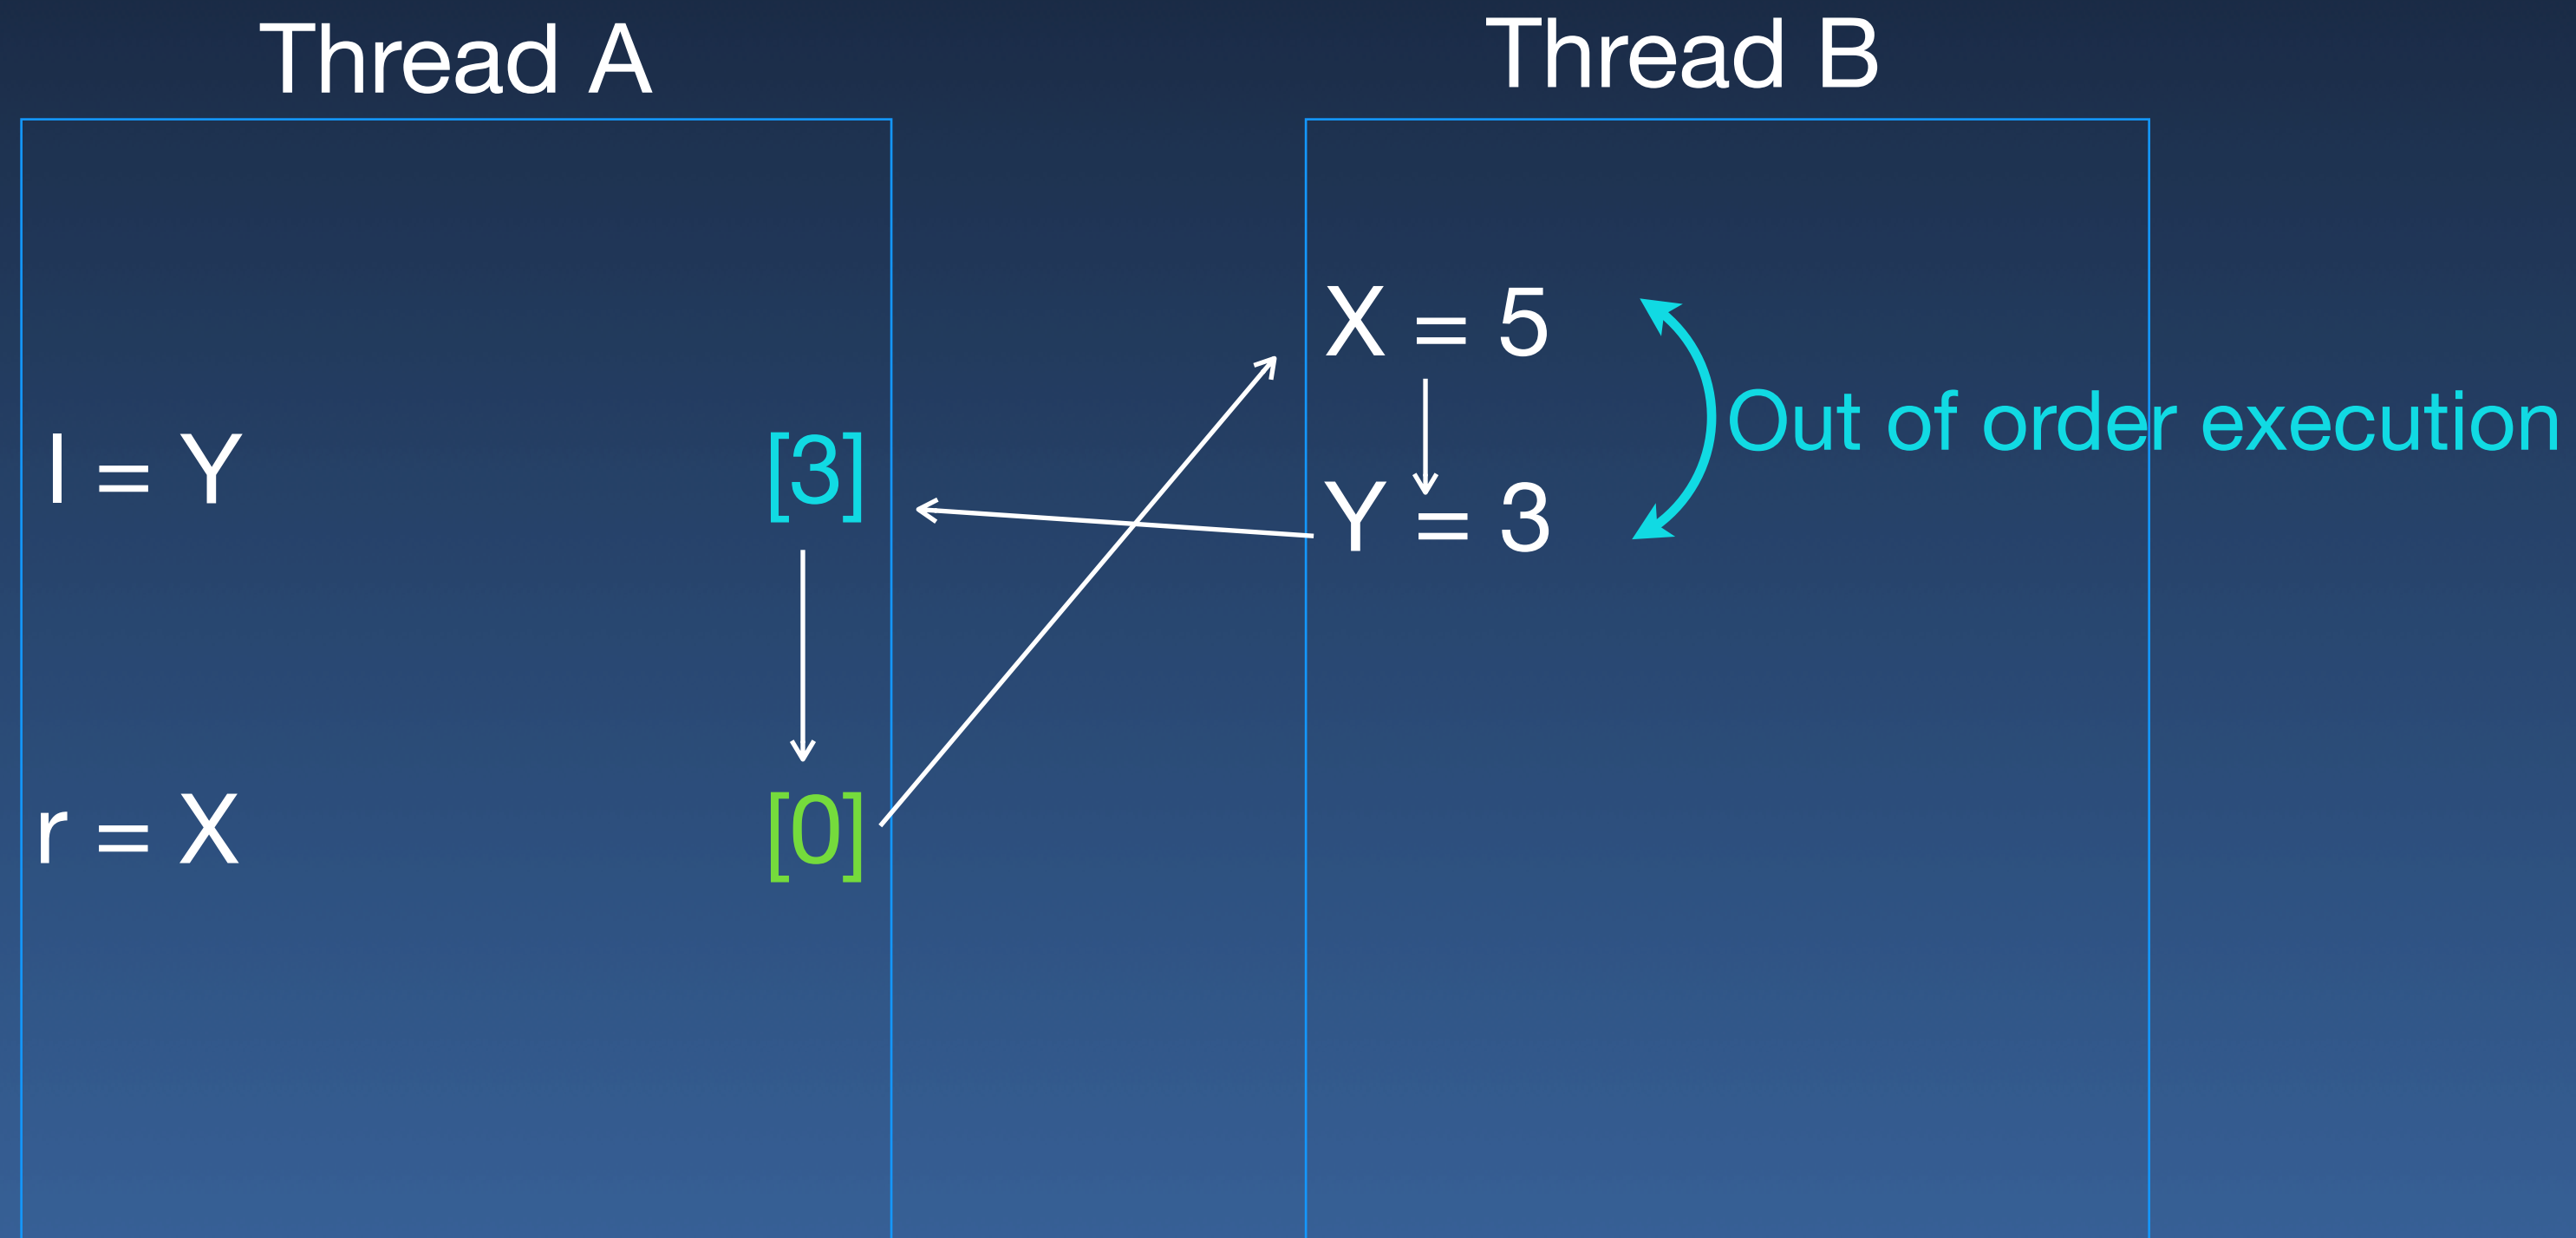
COL380

Introduction to
Parallel & Distributed Programming

- Weak Consistency Models

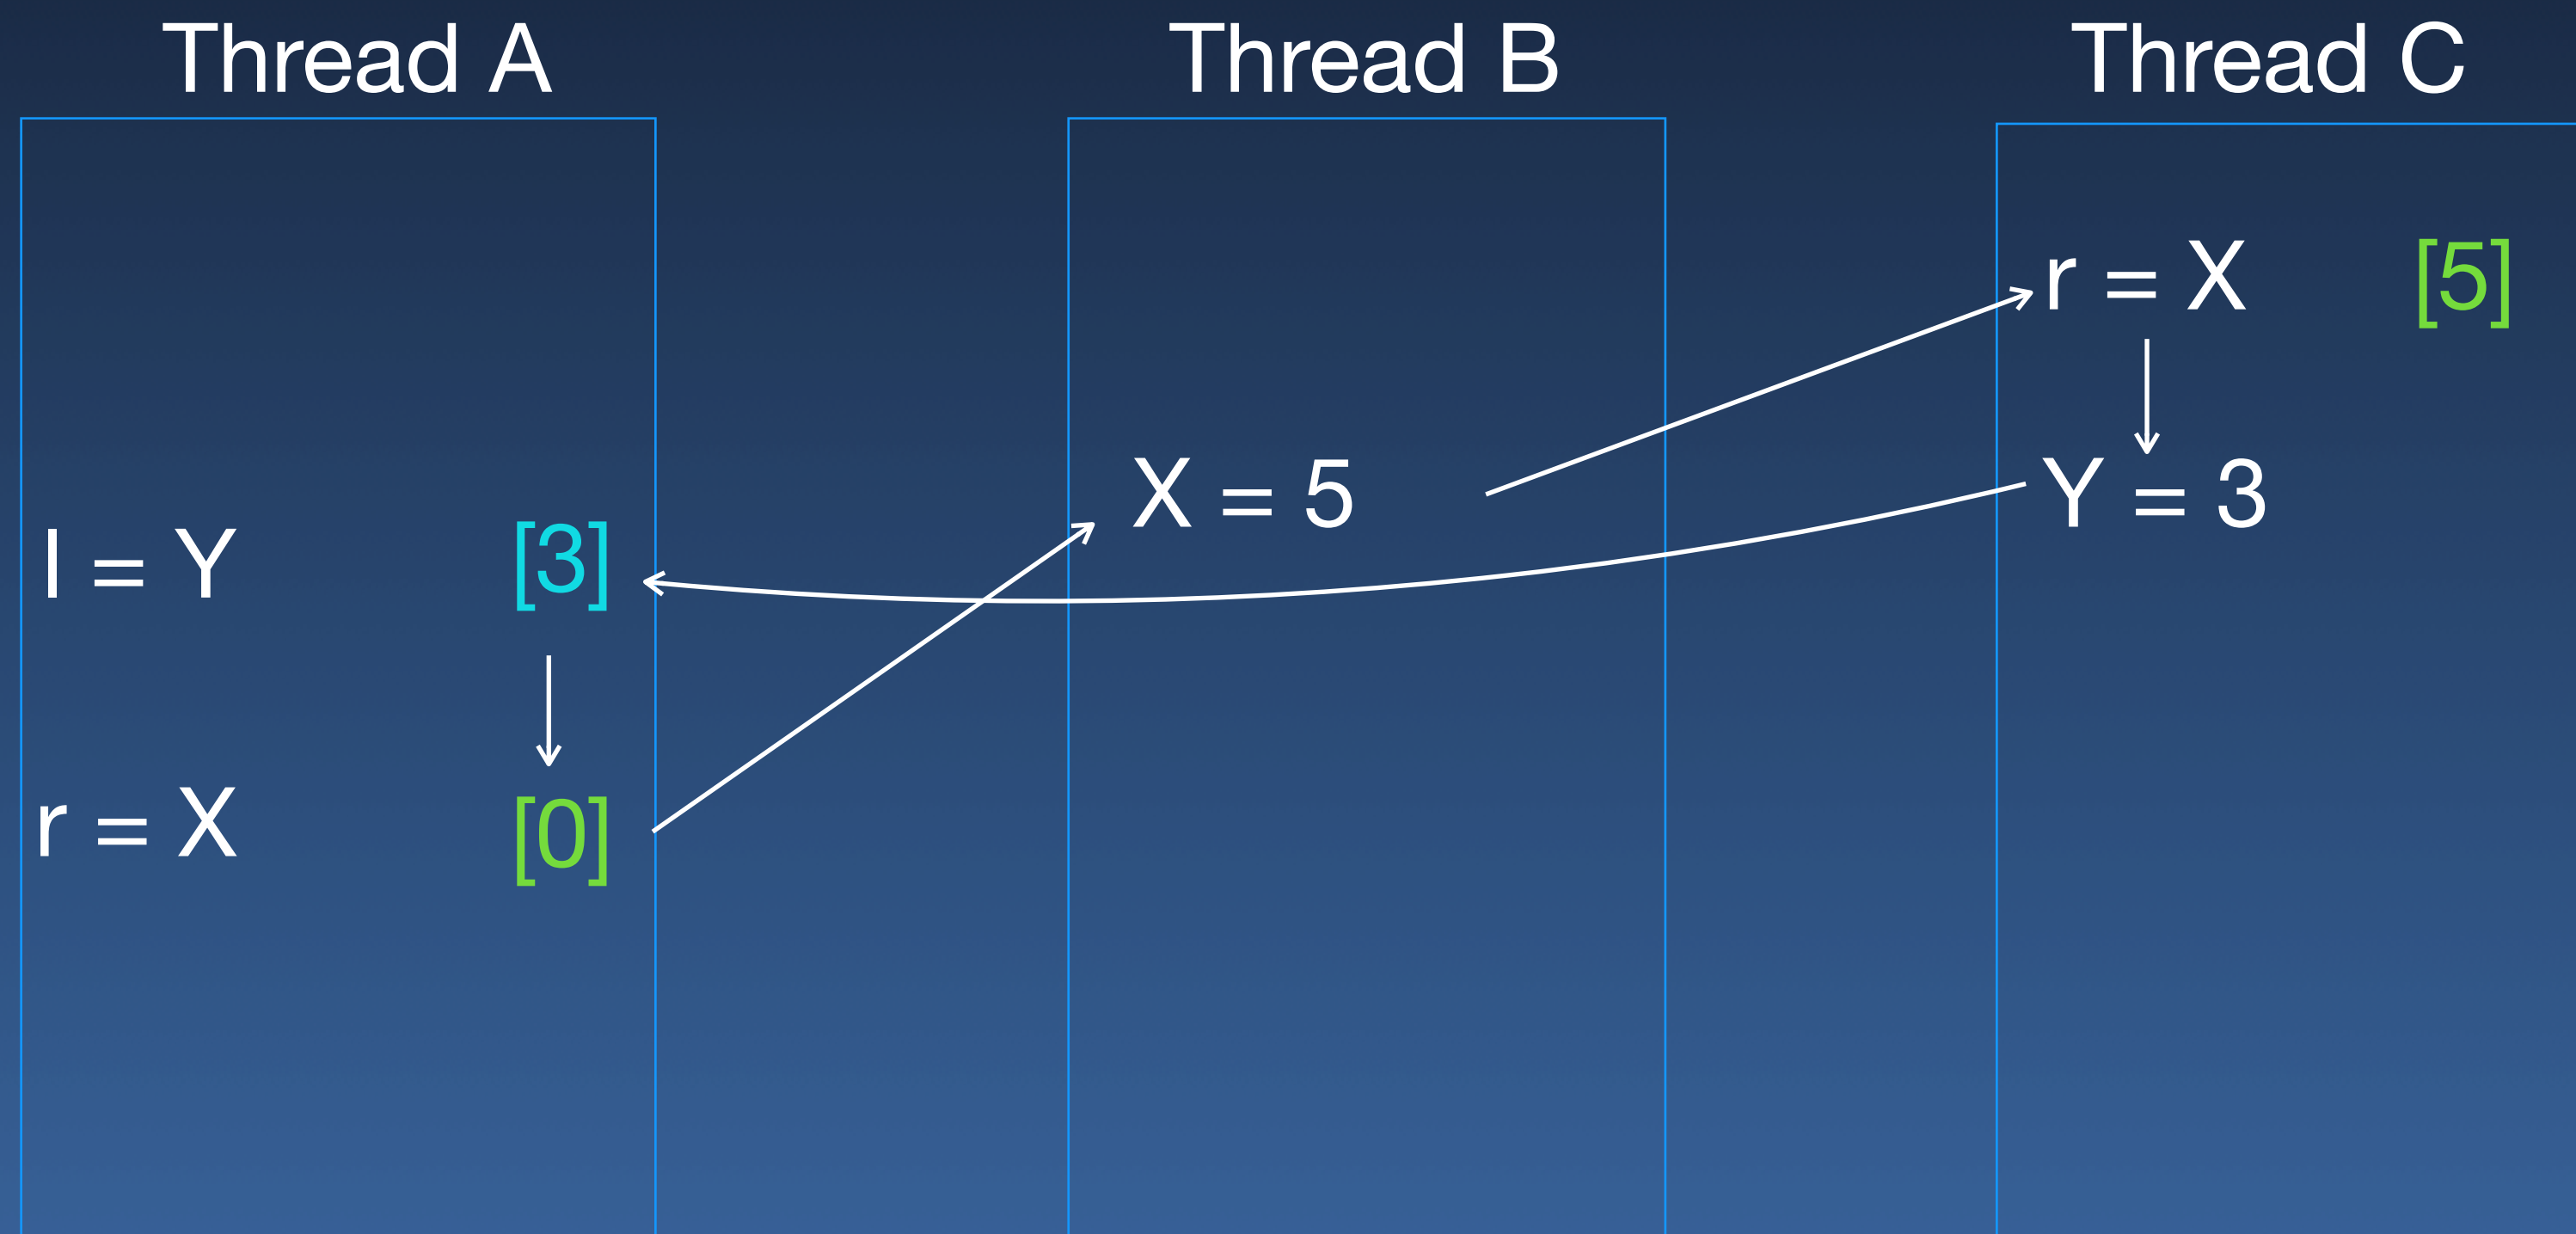
Coherence *vs* Consistency

Initially: $X = 0$; $Y = 0$;



Coherence *vs* Consistency

Initially: $X = 0$; $Y = 0$;



- Consistency is about global state (not per-variable)
 - ➔ Program must not assume higher consistency than available
- Only local memory dependencies visible to compiler/architecture
 - ➔ Can allocate a register or stack entry for some shared variable
 - ➔ Batching of memory transactions
 - ➔ Network can also reorder two memory messages

- Consistency is about global state (not per-variable)
 - Program must not assume higher consistency than available
- Only local memory dependencies visible to compiler/architecture
 - Can allocate a register or stack entry
 - Batching of memory transactions
 - Network can also reorder two memory operations

```
X=1; Y=1; X=2;
```

→ Second write to X may happen before Y's

→ 1st write may never happen

- Consistency is about global state (not per-variable)
 - ➔ Program must not assume higher consistency than available
- Only local memory dependencies visible to compiler/architecture
 - ➔ Can allocate a register or stack entry for some shared variable
 - ➔ Batching of memory transactions
 - ➔ Network can also reorder two memory messages

★ Solutions: Handle inconsistency, force synchronization

Not SC

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
A1 x = a	B1 x = b	C1 y1 = x (b)	D1 z1=x (a)
		C2 y2 = x (a)	D2 z2=x (b)

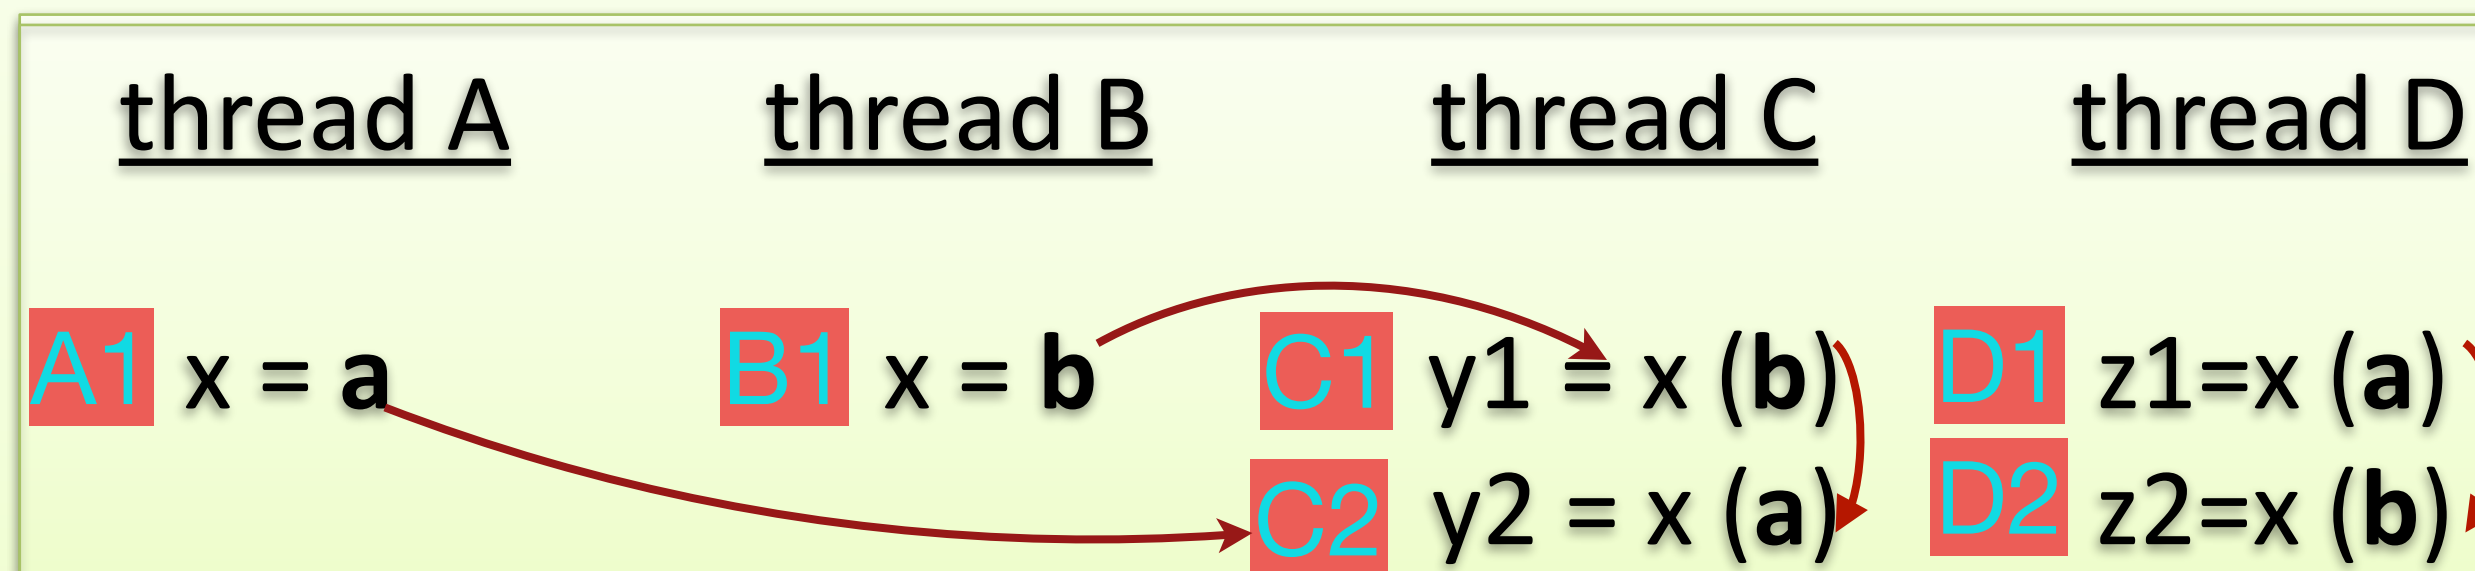
NOT SC

Not SC

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
A1 x = a	B1 x = b	C1 y1 = x (b)	D1 z1=x (a)
		C2 y2 = x (a)	D2 z2=x (b)

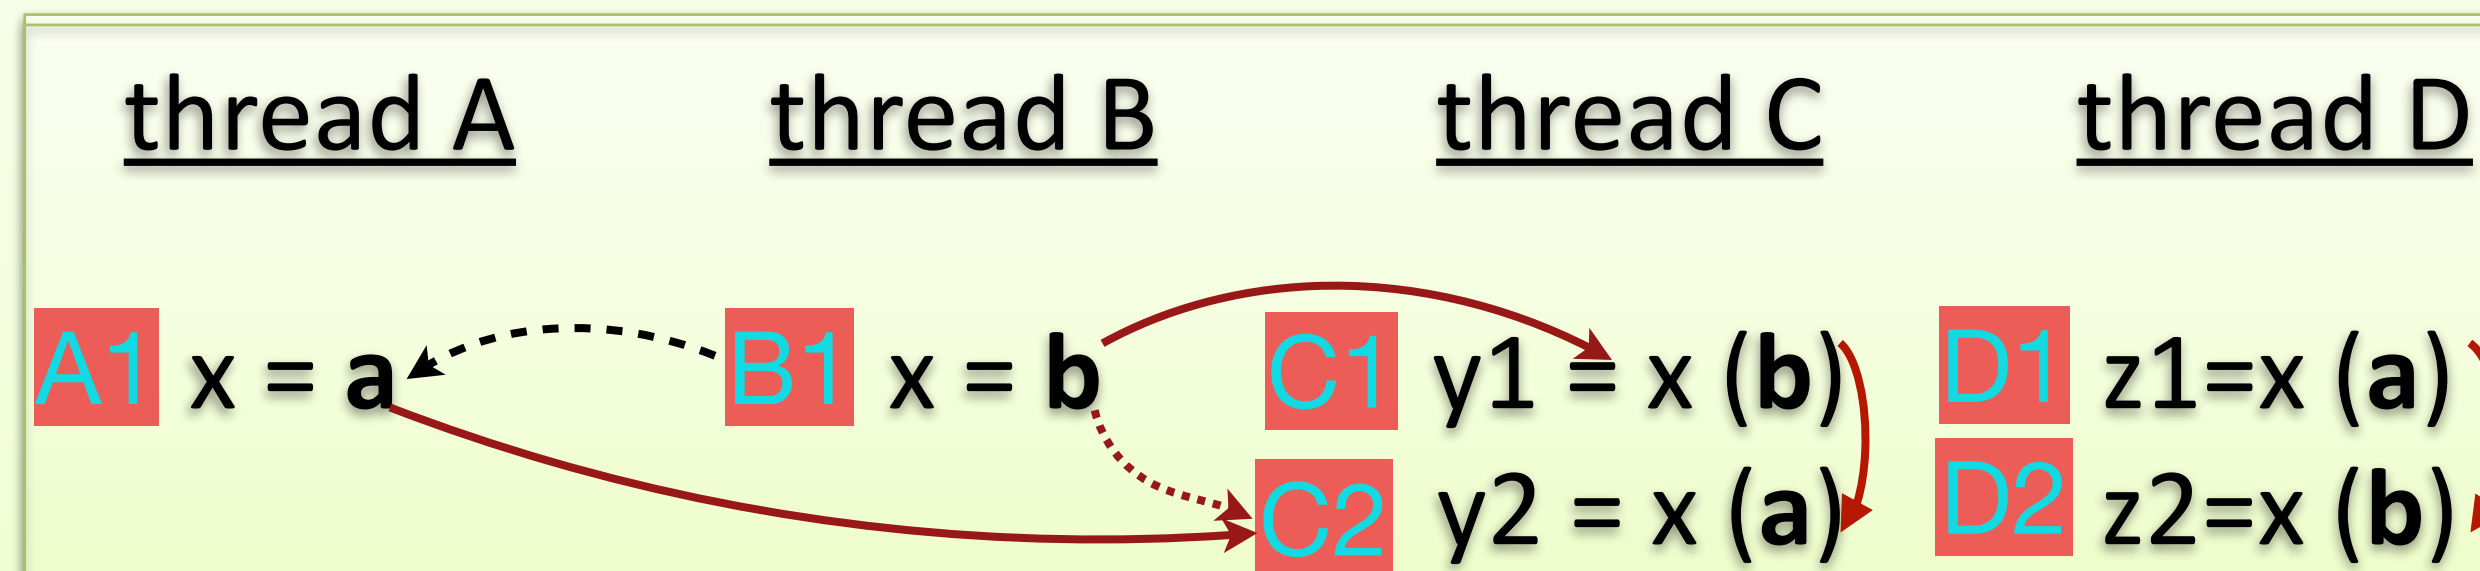
NOT SC

Not SC



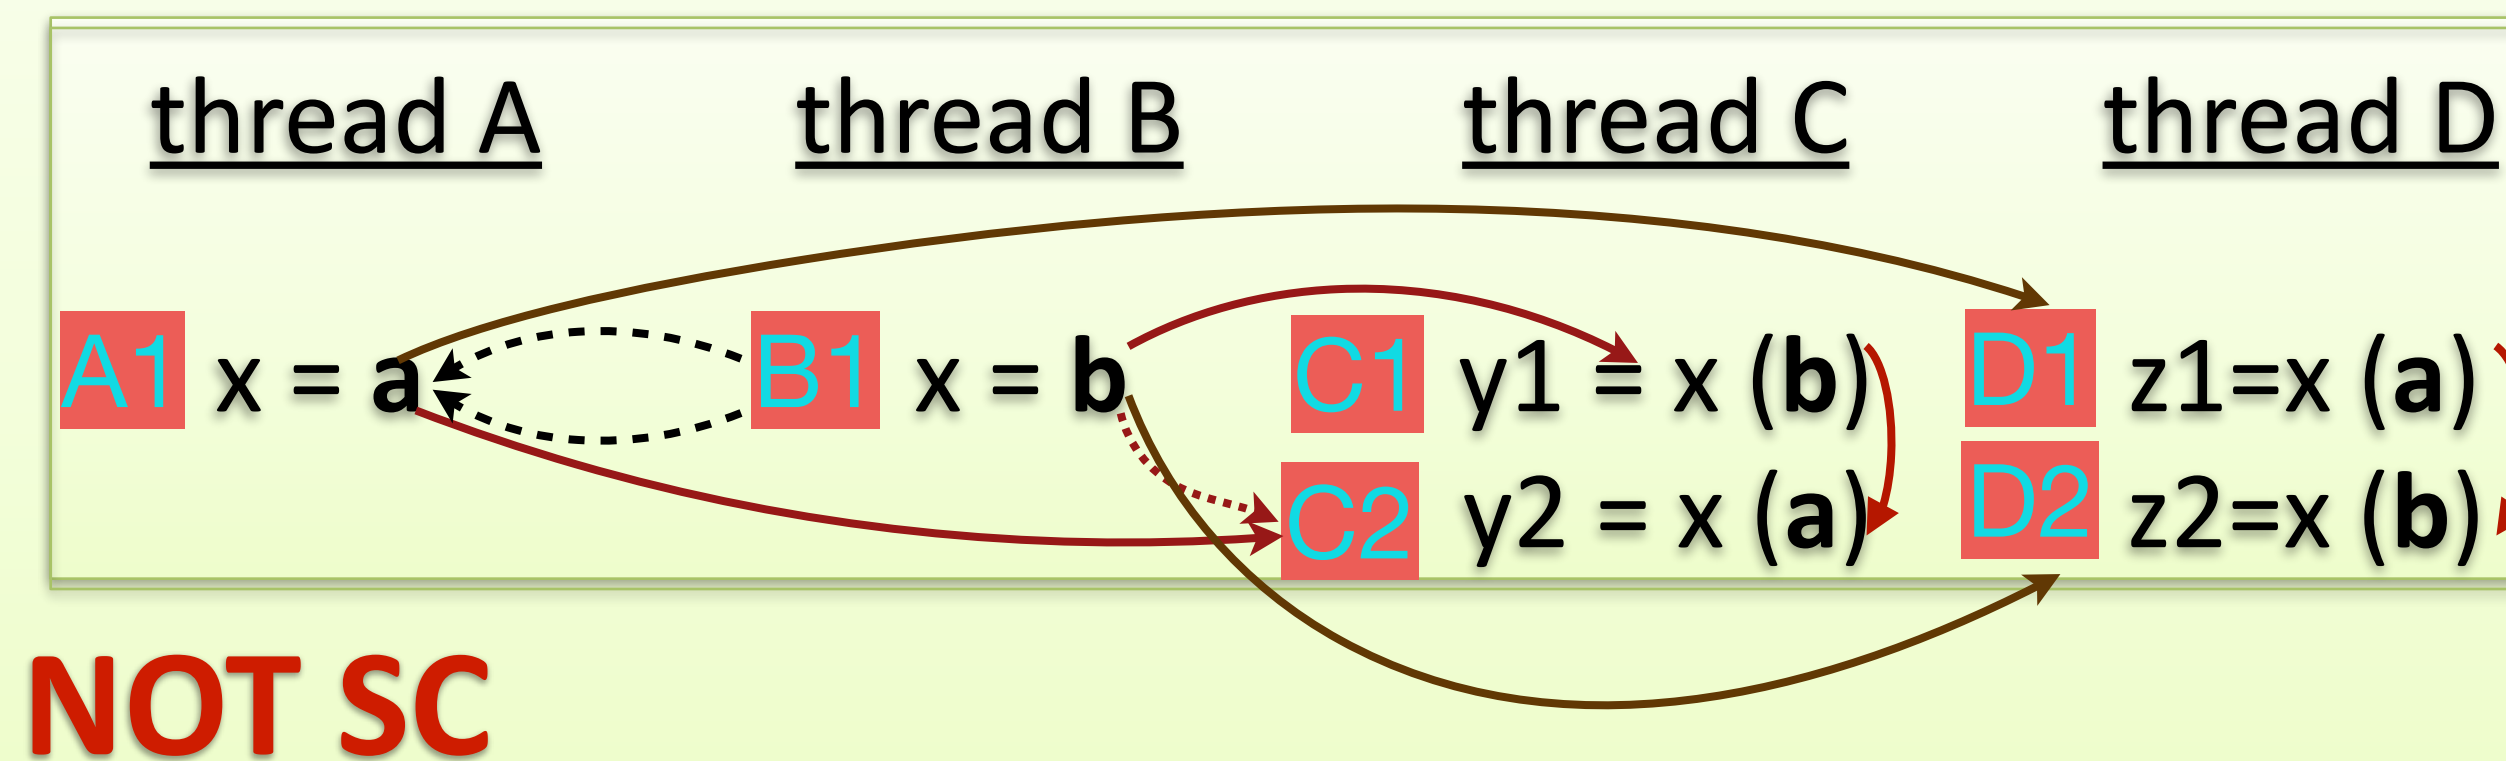
NOT SC

Not SC



NOT SC

Not SC



SC Inefficiency

- Hard to implement efficiently
 - Need to enforce serialization of operations
 - No re-ordering of instructions allowed

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
x = a	x = b	y1 = x (b)	z1=x (a)
		y2 = x (a)	z2=x (b)

- Some solutions:
 - Allow out-of-order execution, Detect and recover from SC violation
 - Only enforce ordering when required
 - ▶ programmer enforced

Relaxing SC

initially: ready=0, data=0

thread1

data1 = 1

data2 = 1

SYNC

thread 2

SYNC

sav1 = data1

sav2 = data2

Relaxing SC

initially: ready=0, data=0

thread1

data1 = 1
data2 = 1
SYNC

thread 2

SYNC
sav1 = data1
sav2 = data2

Causal Consistency

- Write is causally ordered after all earlier reads/writes in its thread
 - ➔ write may depends on the current complete 'state'
- Read is causally ordered after its causative write
- Causality is transitive
- \exists sequential order of causally related operations consistent with every thread's view
 - ➔ Non-related writes may be seen in different order by different threads

Causal Consistency

- Write is causally ordered after all earlier reads/writes in its thread

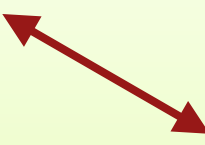
→ write may depends on the current complete 'state'

- Read is causally ordered after its causative write

Causally Consistent

- Causality is transitive

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
x = a		y1 = x (b)	z1=x (a)
	x = b	y2 = x (a)	z2=x (b)



- \exists sequential order of causally related operations consistent with every thread's view

→ Non-related writes may be seen in different order by different threads

Causal Consistency

- Write is causally ordered after all earlier reads/writes in its thread

→ write may depends on the current complete 'state'

- Read is causally ordered after its causative write

~~Causally Consistent~~

- Causality is transitive

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
x = a	y1 = x (a)	y1 = x (b)	z1=x (a)
	x = b	y2 = x (a)	z2=x (b)

- \exists sequential order of causally related operations consistent with every thread's view

→ Non-related writes may be seen in different order by different threads

- All threads see **all writes** by each thread in the order of that thread
 - ➔ all instances of write(**x**) are seen by each thread in the same order
 - ➔ No need to consistently order writes to different variables by different threads
- Easy to implement
 - ➔ Two or more writes from a single source must remain in order, as in a pipeline
 - ➔ All writes are through to the memory

Processor Consistency

- All threads see **all writes** by each thread in the order of that thread
 - ➔ all instances of write(**x**) are seen by each thread in **FIFO consistency** relaxes this constraint
 - ➔ No need to consistently order writes to different variables by different threads
- Easy to implement
 - ➔ Two or more writes from a single source must remain in order, as in a pipeline
 - ➔ All writes are through to the memory

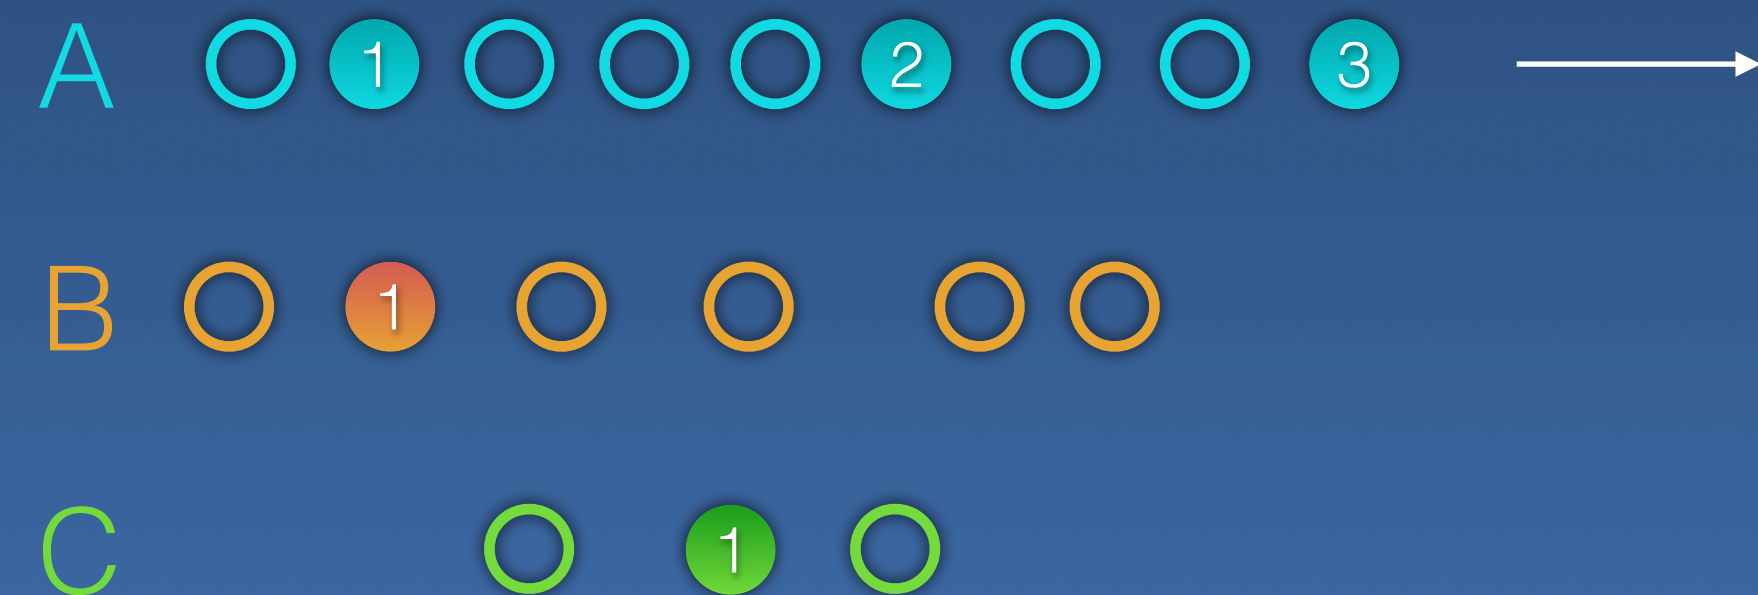
Processor Consistency

- All threads see **all writes** by each thread in the order of that thread
 - ➔ all instances of write(**x**) are seen by each thread in the same order
 - ➔ No need to consistently order writes to different variables by different threads
- Easy to implement
 - ➔ Two or more writes from a single source must remain in order, as in a pipeline
 - ➔ All writes are through to the memory

FIFO consistency is also known as **PRAM consistency**

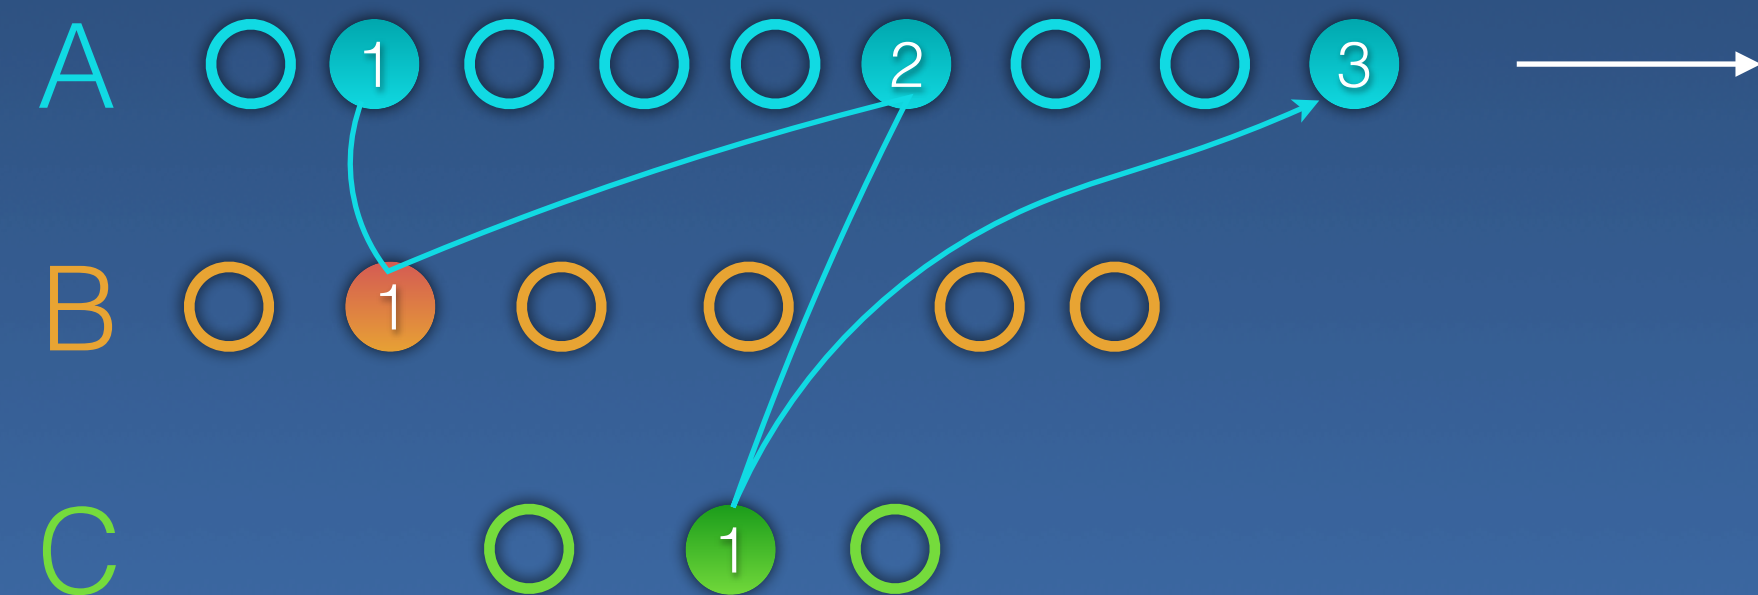
Weak Consistency

- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
 - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
 - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations



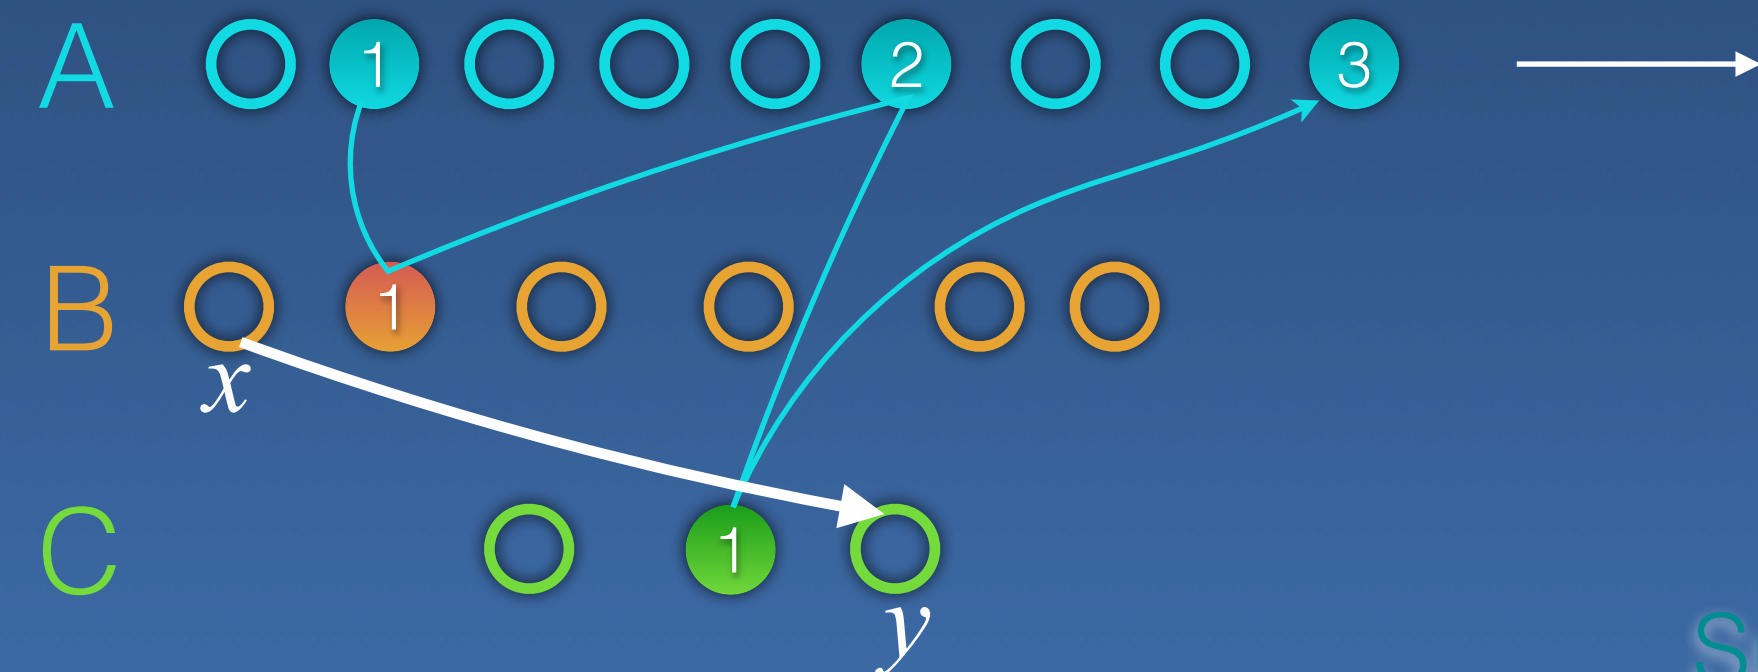
Weak Consistency

- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
 - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
 - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations



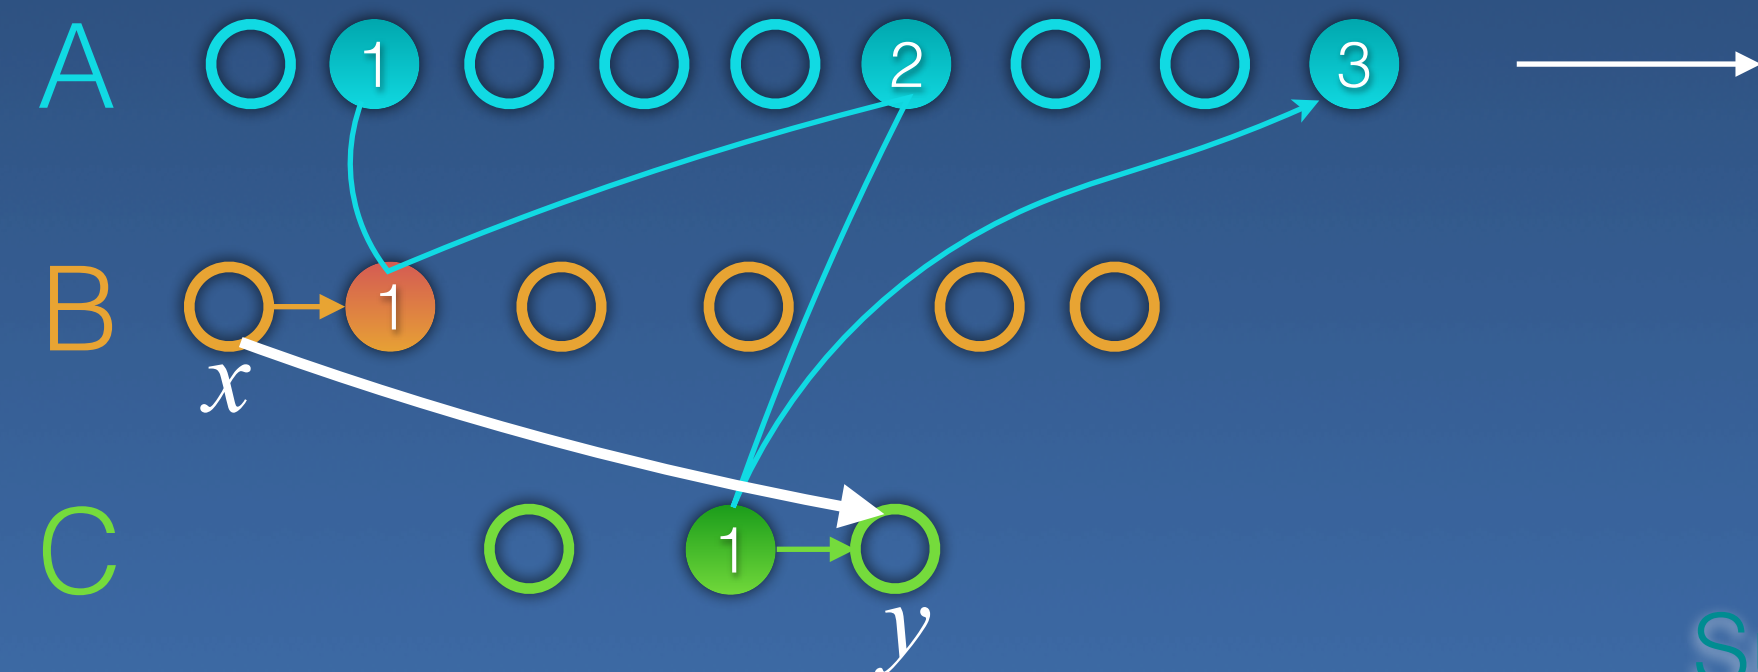
Weak Consistency

- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
 - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
 - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations



Weak Consistency

- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
 - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
 - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations



OpenMP Flush

flagA = flagB = 0

Thread A

flagA = 1;

```
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

flagB = 1;

```
if (flagA == 0) {  
    shared++;  
}
```

OpenMP Flush

flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush  
if (flagA == 0) {  
    shared++;  
}
```


OpenMP Flush

flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush  
if (flagA == 0) {  
    shared++;  
}
```

Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

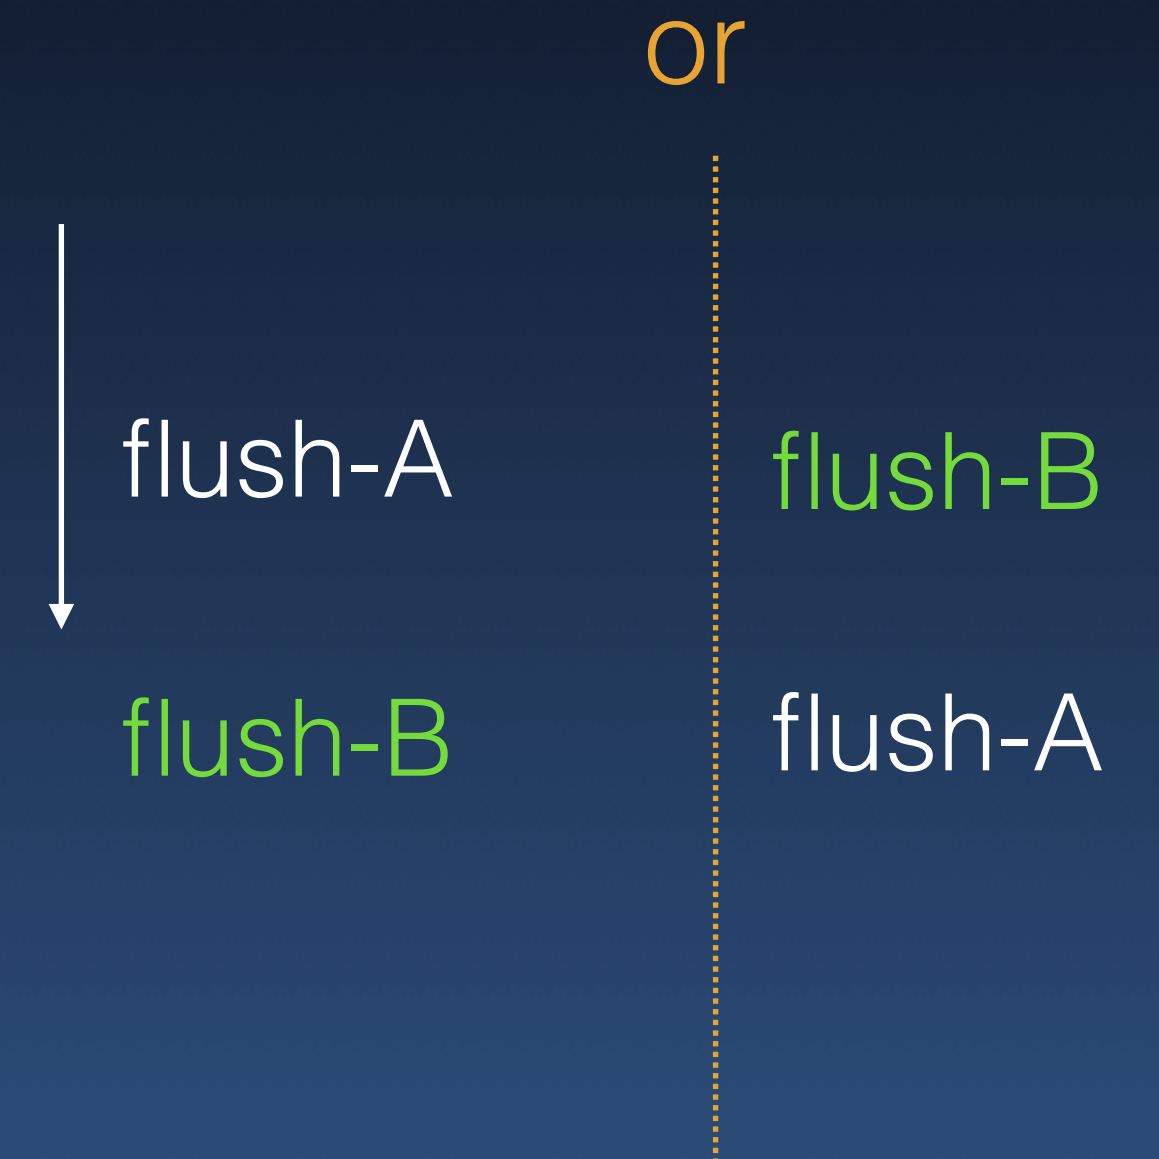
flagA = flagB = 0

Thread A

```
flagA = 1;
#pragma omp flush
if (flagB == 0) {
    shared ++;
}
```

Thread B

```
flagB = 1;
#pragma omp flush
if (flagA == 0) {
    shared++;
}
```



Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

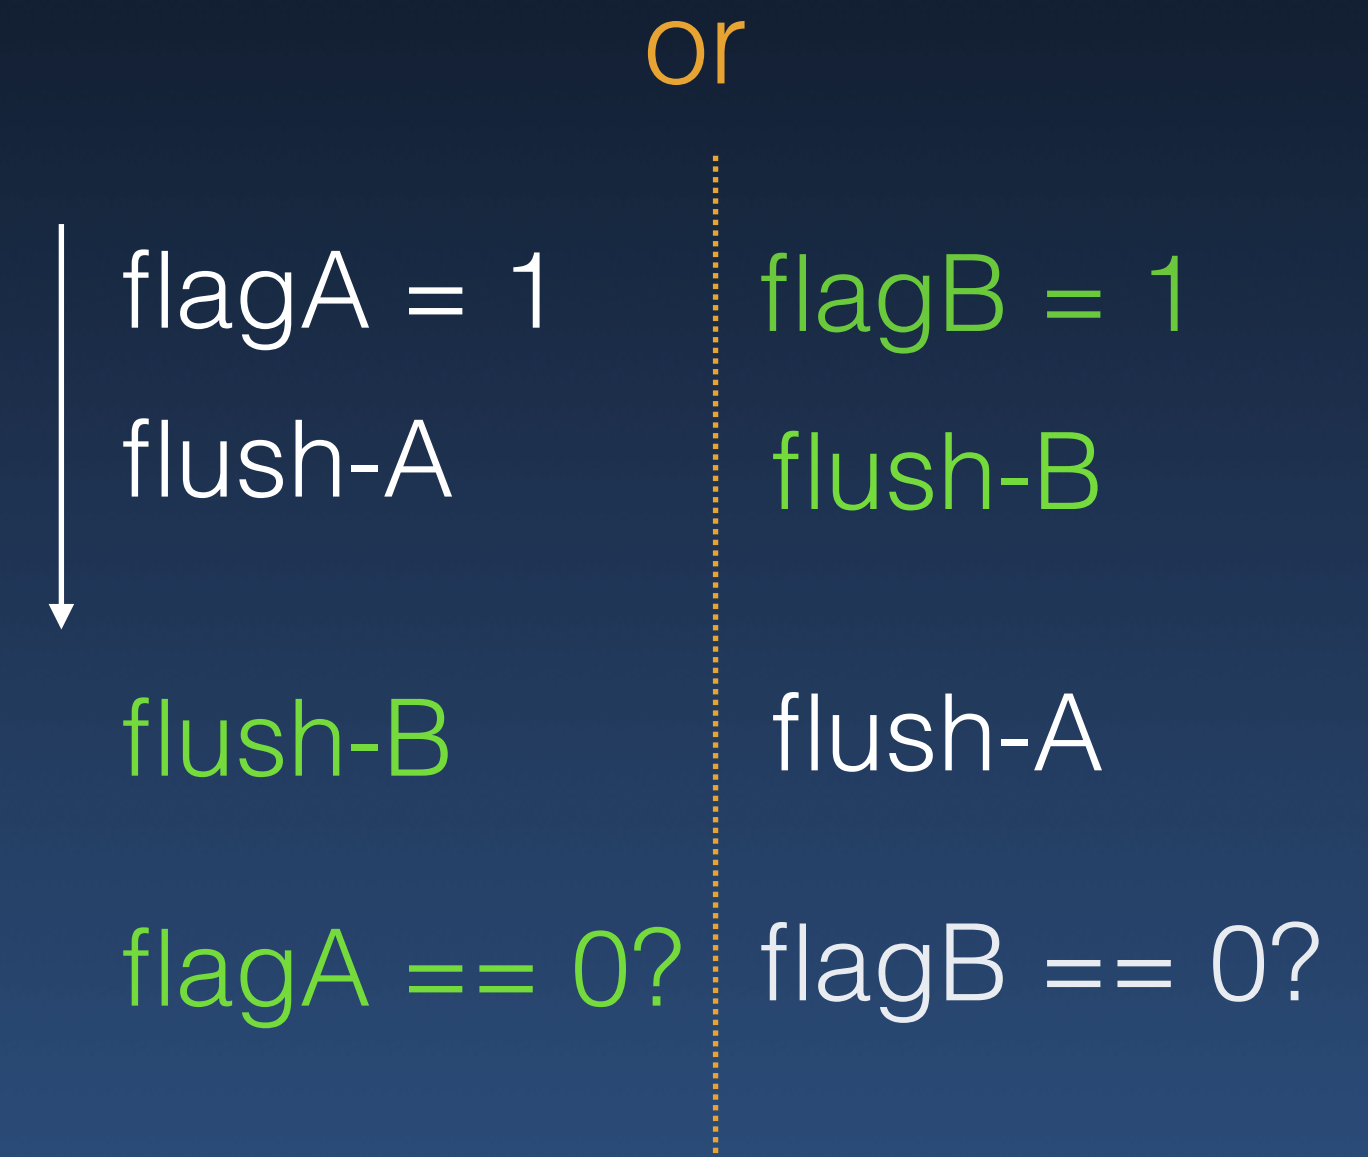
flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush  
if (flagA == 0) {  
    shared ++;  
}
```



Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

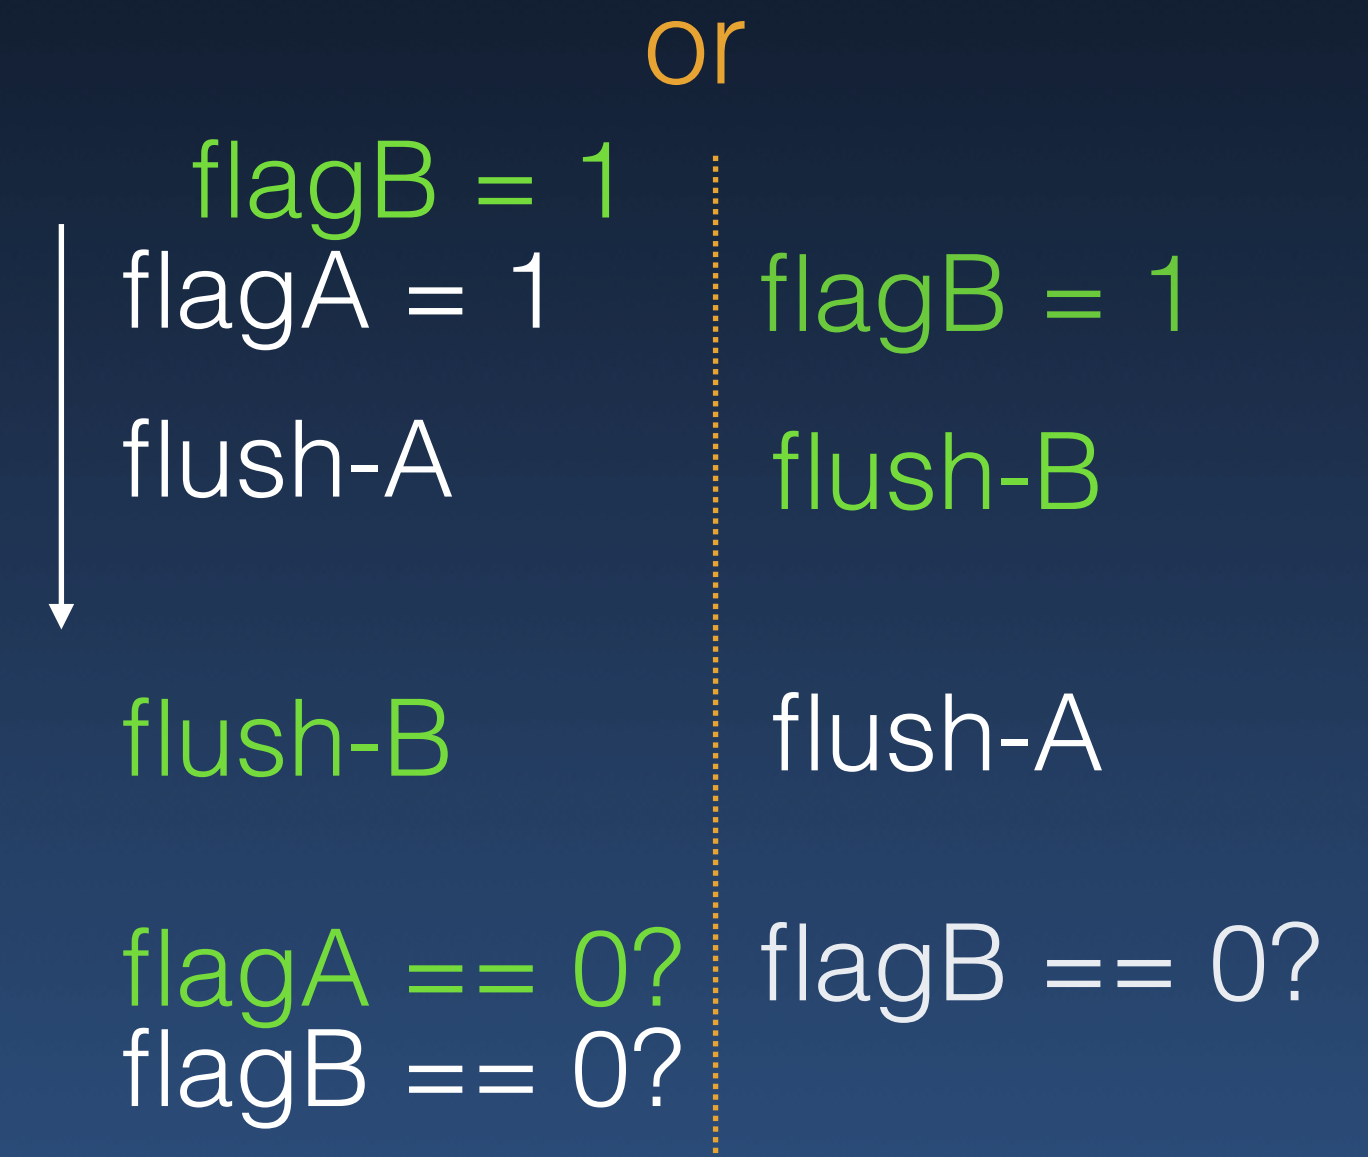
flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush  
if (flagA == 0) {  
    shared ++;  
}
```



Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

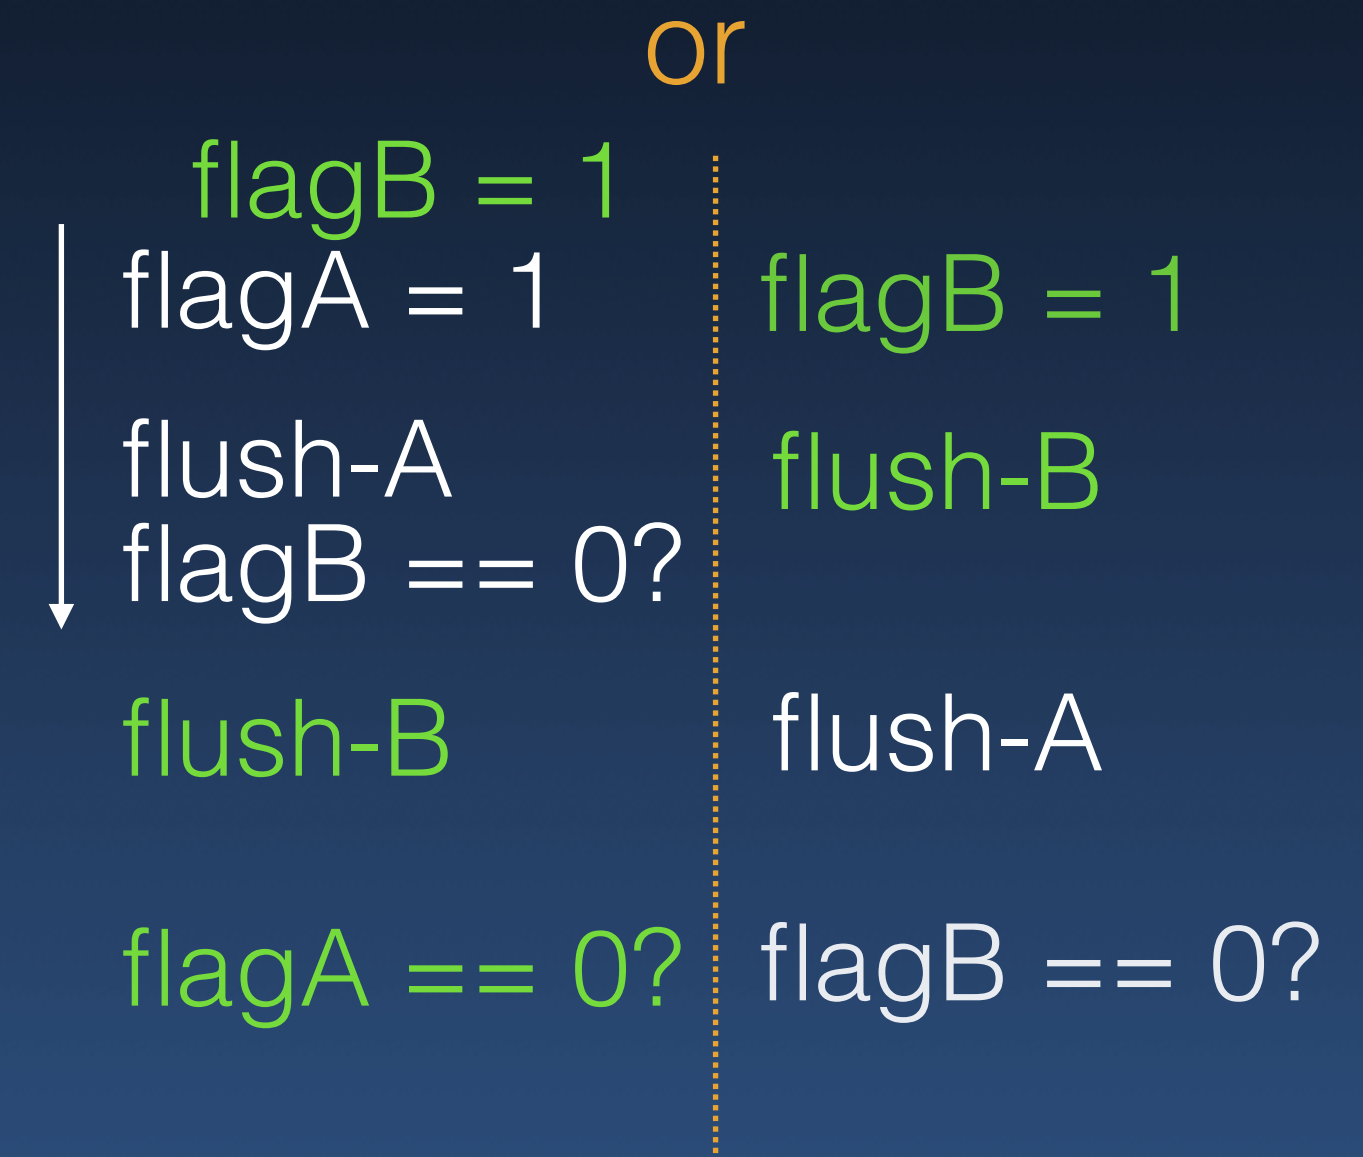
flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush  
if (flagA == 0) {  
    shared ++;  
}
```



Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

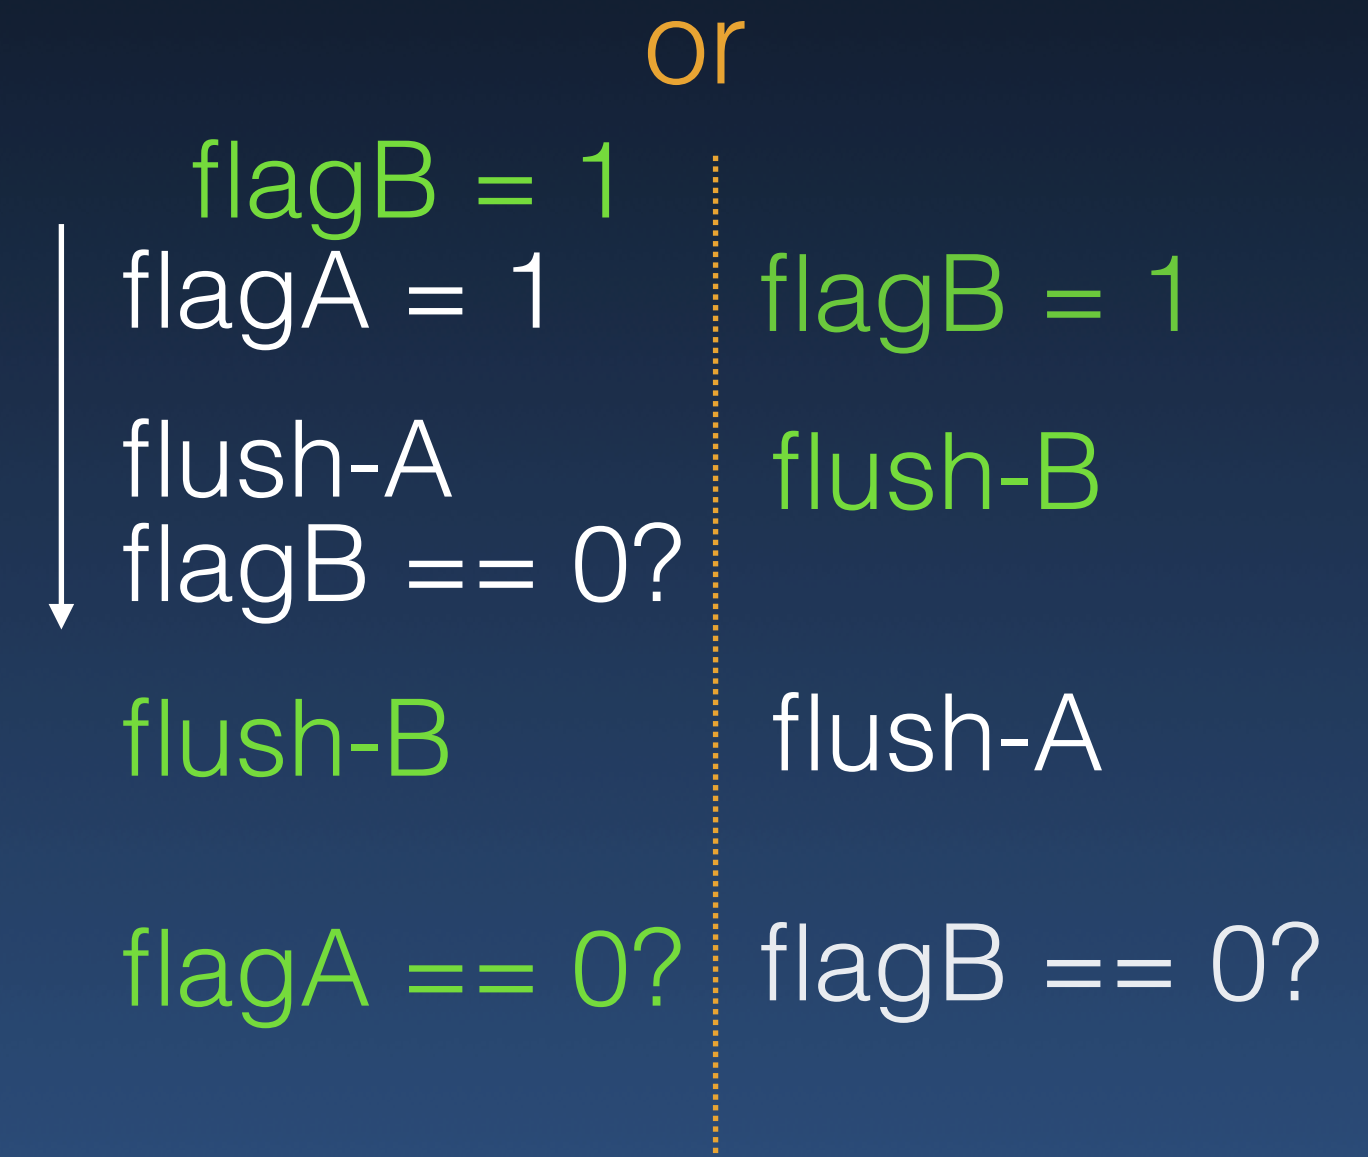
flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush  
if (flagB == 0) {  
    shared ++;  
}  
flagA = 0;  
#pragma omp flush
```

Thread B

```
flagB = 1;  
#pragma omp flush  
if (flagA == 0) {  
    shared ++;  
}  
flagB = 0;  
#pragma omp flush
```



Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush (flagA, flagB)  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush (flagA, flagB)  
if (flagA == 0) {  
    shared++;  
}
```

Access other variables

Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush (flagA, flagB)  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush (flagA, flagB)  
if (flagA == 0) {  
    shared ++;  
}
```

Access other variables



Flush (Memory fence) are sequentially consistent
Wait for ongoing memory operations to finish
Discard Local view (Subsequent reads actually load from memory)

OpenMP Flush

flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush (flagA, flagB)  
if (flagB == 0) {  
    shared ++;  
}
```

Access other variables

Thread B

```
flagB = 1;  
#pragma omp flush (flagA, flagB)  
if (flagA == 0) {  
    shared ++;  
}
```

Implicit Flush for all
synchronization operations

```
#pragma omp atomic read  
val = var;  
#pragma omp atomic write  
var = expr();
```

Flush (Memory fence) are sequentially consistent

Wait for ongoing memory operations to finish

Discard Local view (Subsequent reads actually load from memory)


OpenMP Flush

Operations before **Release** flush must appear before (Completes)
Operations after **Acquire** flush must appear after (Initiates)

flagA = flagB = 0


Thread A

```
flagA = 1;
#pragma omp flush release
if (flagB == 0) {
    shared++;
}
```



Thread B

```
flagB = 1;
#pragma omp flush acquire
if (flagA == 0) {
    shared++;
}
```



flagA = 1
flush-A

flush-B

flagA == 0?

or

flagB = 1
flush-B

flush-A

flagB == 0?

Flush (Memory fence) are sequentially consistent
Wait for ongoing memory operations to finish
Discard Local view (Subsequent reads actually load from memory)

Operations before **Release** flush must appear before (Completes)
Operations after **Acquire** flush must appear after (Initiates)

OpenMP Flush

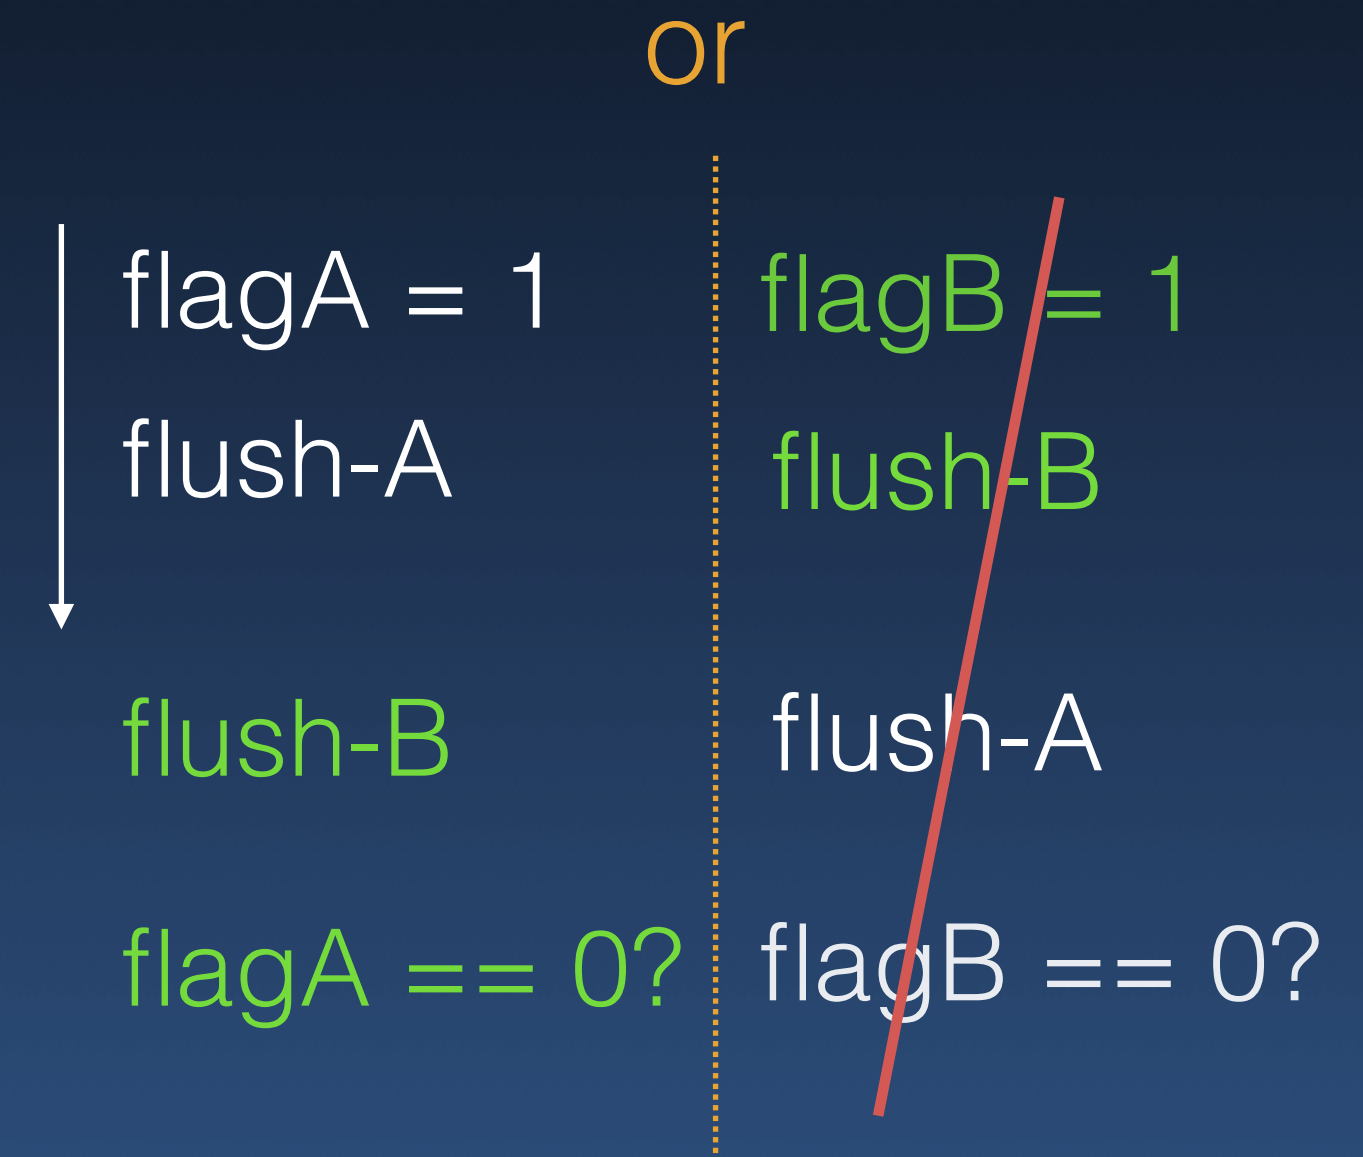
flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush release  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush acquire  
if (flagA == 0) {  
    shared ++;  
}
```



Flush (Memory fence) are sequentially consistent
Wait for ongoing memory operations to finish
Discard Local view (Subsequent reads actually load from memory)

Operations before **Release** flush must appear before (Completes)
Operations after **Acquire** flush must appear after (Initiates)

OpenMP Flush

flagA = flagB = 0

Thread A

```
flagA = 1;  
#pragma omp flush release  
if (flagB == 0) {  
    shared ++;  
}
```

Thread B

```
flagB = 1;  
#pragma omp flush acquire  
if (flagA == 0) {  
    shared ++;  
}
```

or

flagA = 1	flagB = 1
flush-A	flush-B
flush-B	flush-A
flagA == 0?	flagB == 0?

Flush (Memory fence) are sequentially consistent
Wait for ongoing memory operations to finish
Discard Local view (Subsequent reads actually load from memory)

Producer-Consumer

```
int data, flag = 0;
```

Thread 0

// Produce data

`data = 42;`

// Set flag to signal Thread 1

`flag = 1;`

Thread 1

// Busy-wait until flag is signalled

`while (flag != 1) {`

`}`

// Consume data

`printf("data=%d\n", data);`

Producer-Consumer

```
int data, flag = 0;
```

Thread 0

// Produce data

```
data = 42;
```

// Set flag to signal Thread 1

```
flag = 1;
```

Thread 1

// Busy-wait until flag is signalled

```
while (flag != 1) {  
    }  
}
```

// Consume data

```
printf("data=%d\n", data);
```

Producer-Consumer

```
int data, flag = 0;
```

Thread 0

// Produce data

`data = 42;`

// Set flag to signal Thread 1

`flag = 1;`

Thread 1

// Busy-wait until flag is signalled

`while (flag != 1) {`

`}`

// Consume data

`printf("data=%d\n", data);`



Producer-Consumer

```
int data, flag = 0;
```

Thread 0

// Produce data

`data = 42;`

// Set flag to signal Thread 1

`flag = 1;`

// Flush

`#pragma omp flush(flag)`

Thread 1

// Busy-wait until flag is signalled

`#pragma omp flush(flag)`

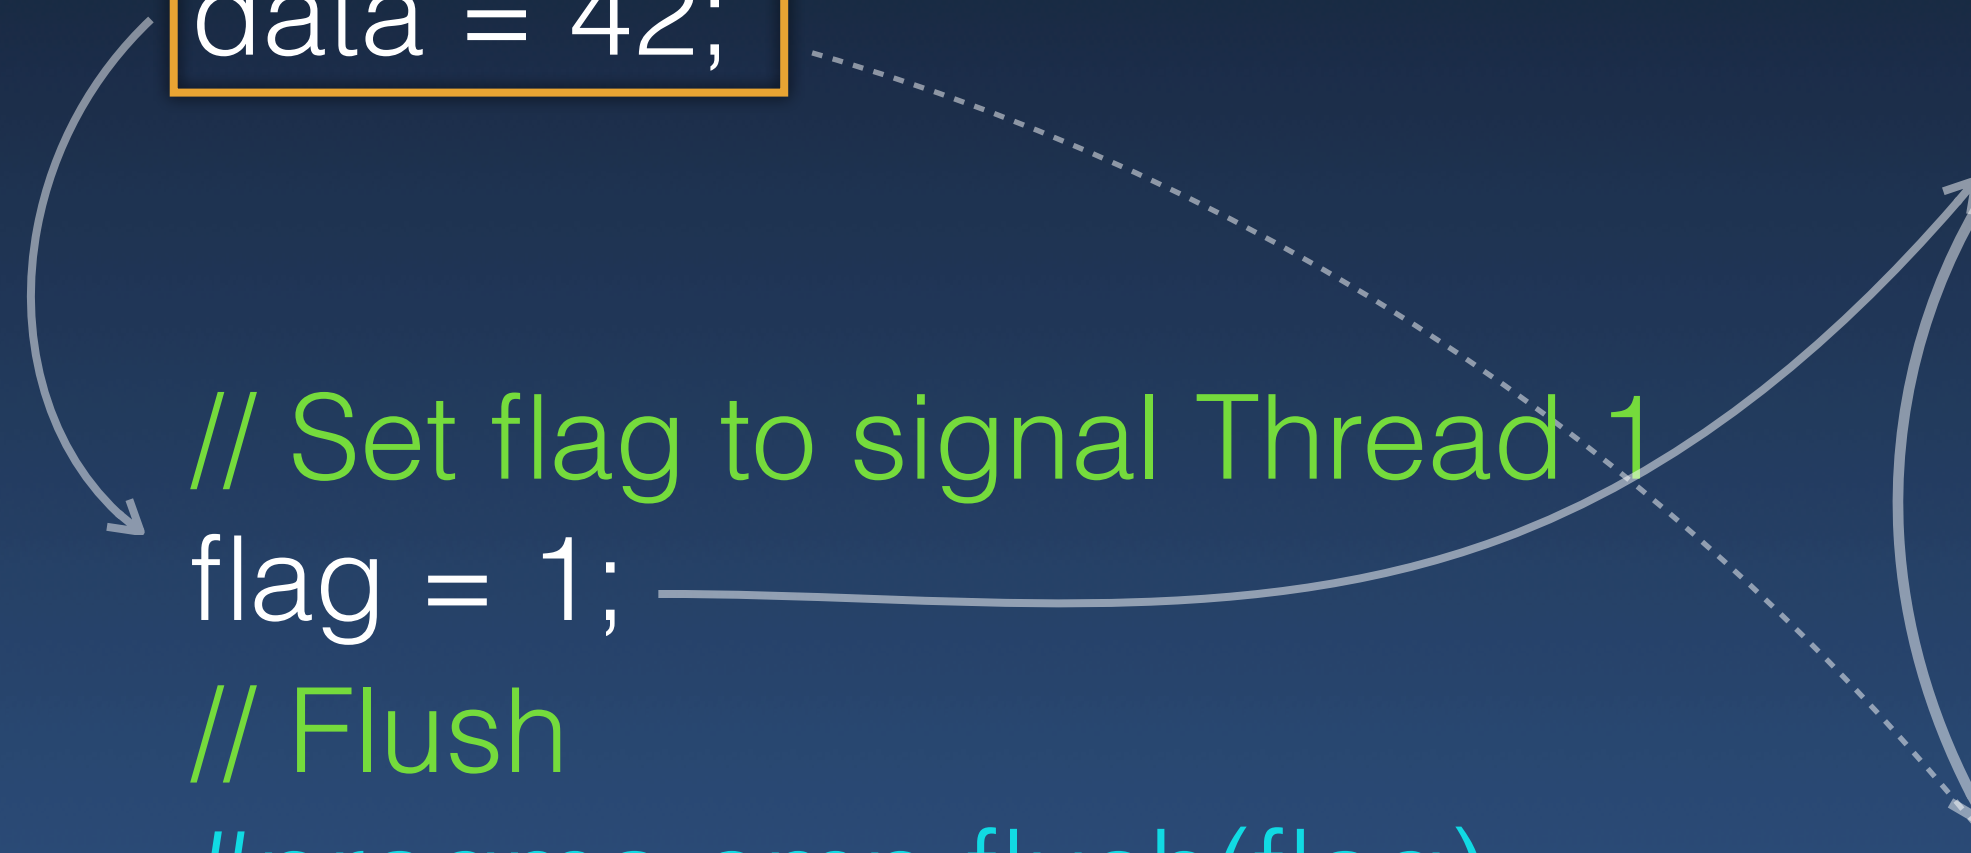
`while (flag != 1) {`

`#pragma omp flush(flag)`

`}`

// Consume data

`printf("data=%d\n", data);`



Producer-Consumer

```
int data, flag = 0;
```

Thread 0

// Produce data

`data = 42;`

// Set flag to signal Thread 1

`flag = 1;`

// Flush

`#pragma omp flush(flag)`

Thread 1

// Busy-wait until flag is signalled

`#pragma omp flush(flag)` consume

`while (flag != 1) {`

`#pragma omp flush(flag)`

`}`

// Consume data

`printf("data=%d\n", data);`



Producer-Consumer

int data, flag = 0;

Thread 0

// Produce data

data = 42;

// Set flag to signal Thread 1

flag = 1;

// Flush

#pragma omp flush(flag)

produce

Thread 1

// Busy-wait until flag is signalled

#pragma omp flush(flag)

consume

while (flag != 1) {

#pragma omp flush(flag)

}

// Consume data

printf(data="%d\n", data);



Producer-Consumer

```
int data, flag = 0;
```

Thread 0

// Produce data

`data = 42;`

// Flush

`#pragma omp flush(flag, data)`

// Set flag to signal Thread 1

`flag = 1;`

// Flush

`#pragma omp flush(flag)`

Thread 1

// Busy-wait until flag is signalled

`#pragma omp flush(flag)`

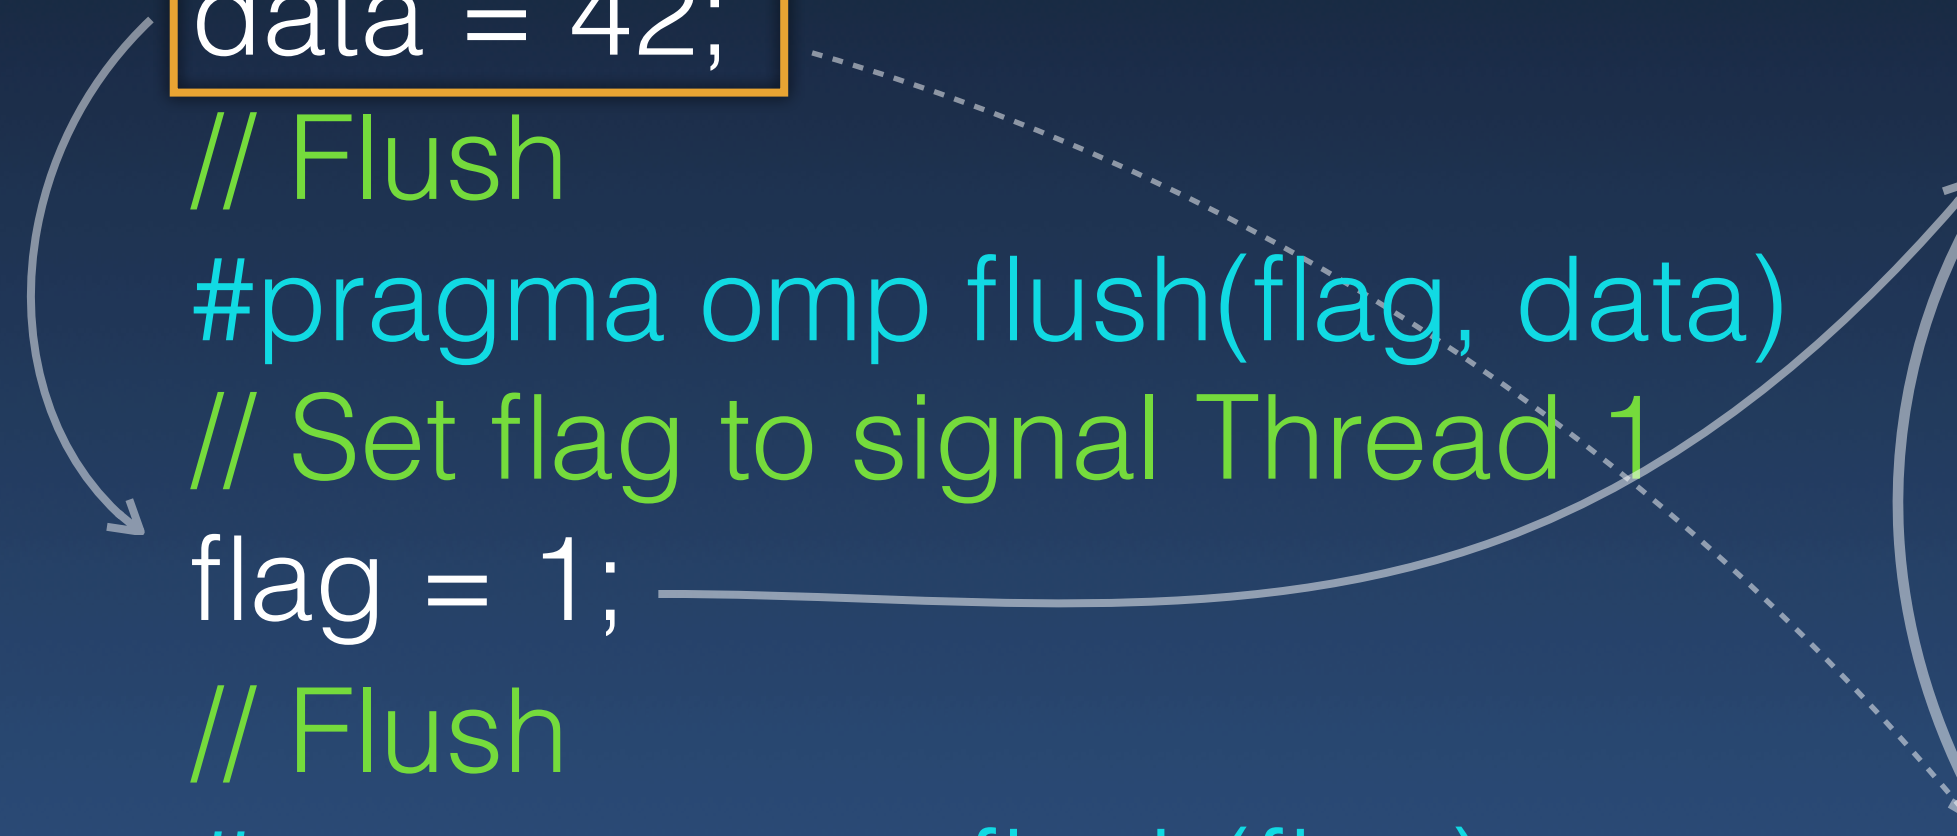
`while (flag != 1) {`

`#pragma omp flush(flag, data)`

`}`

// Consume data

`printf("data=%d\n", data);`

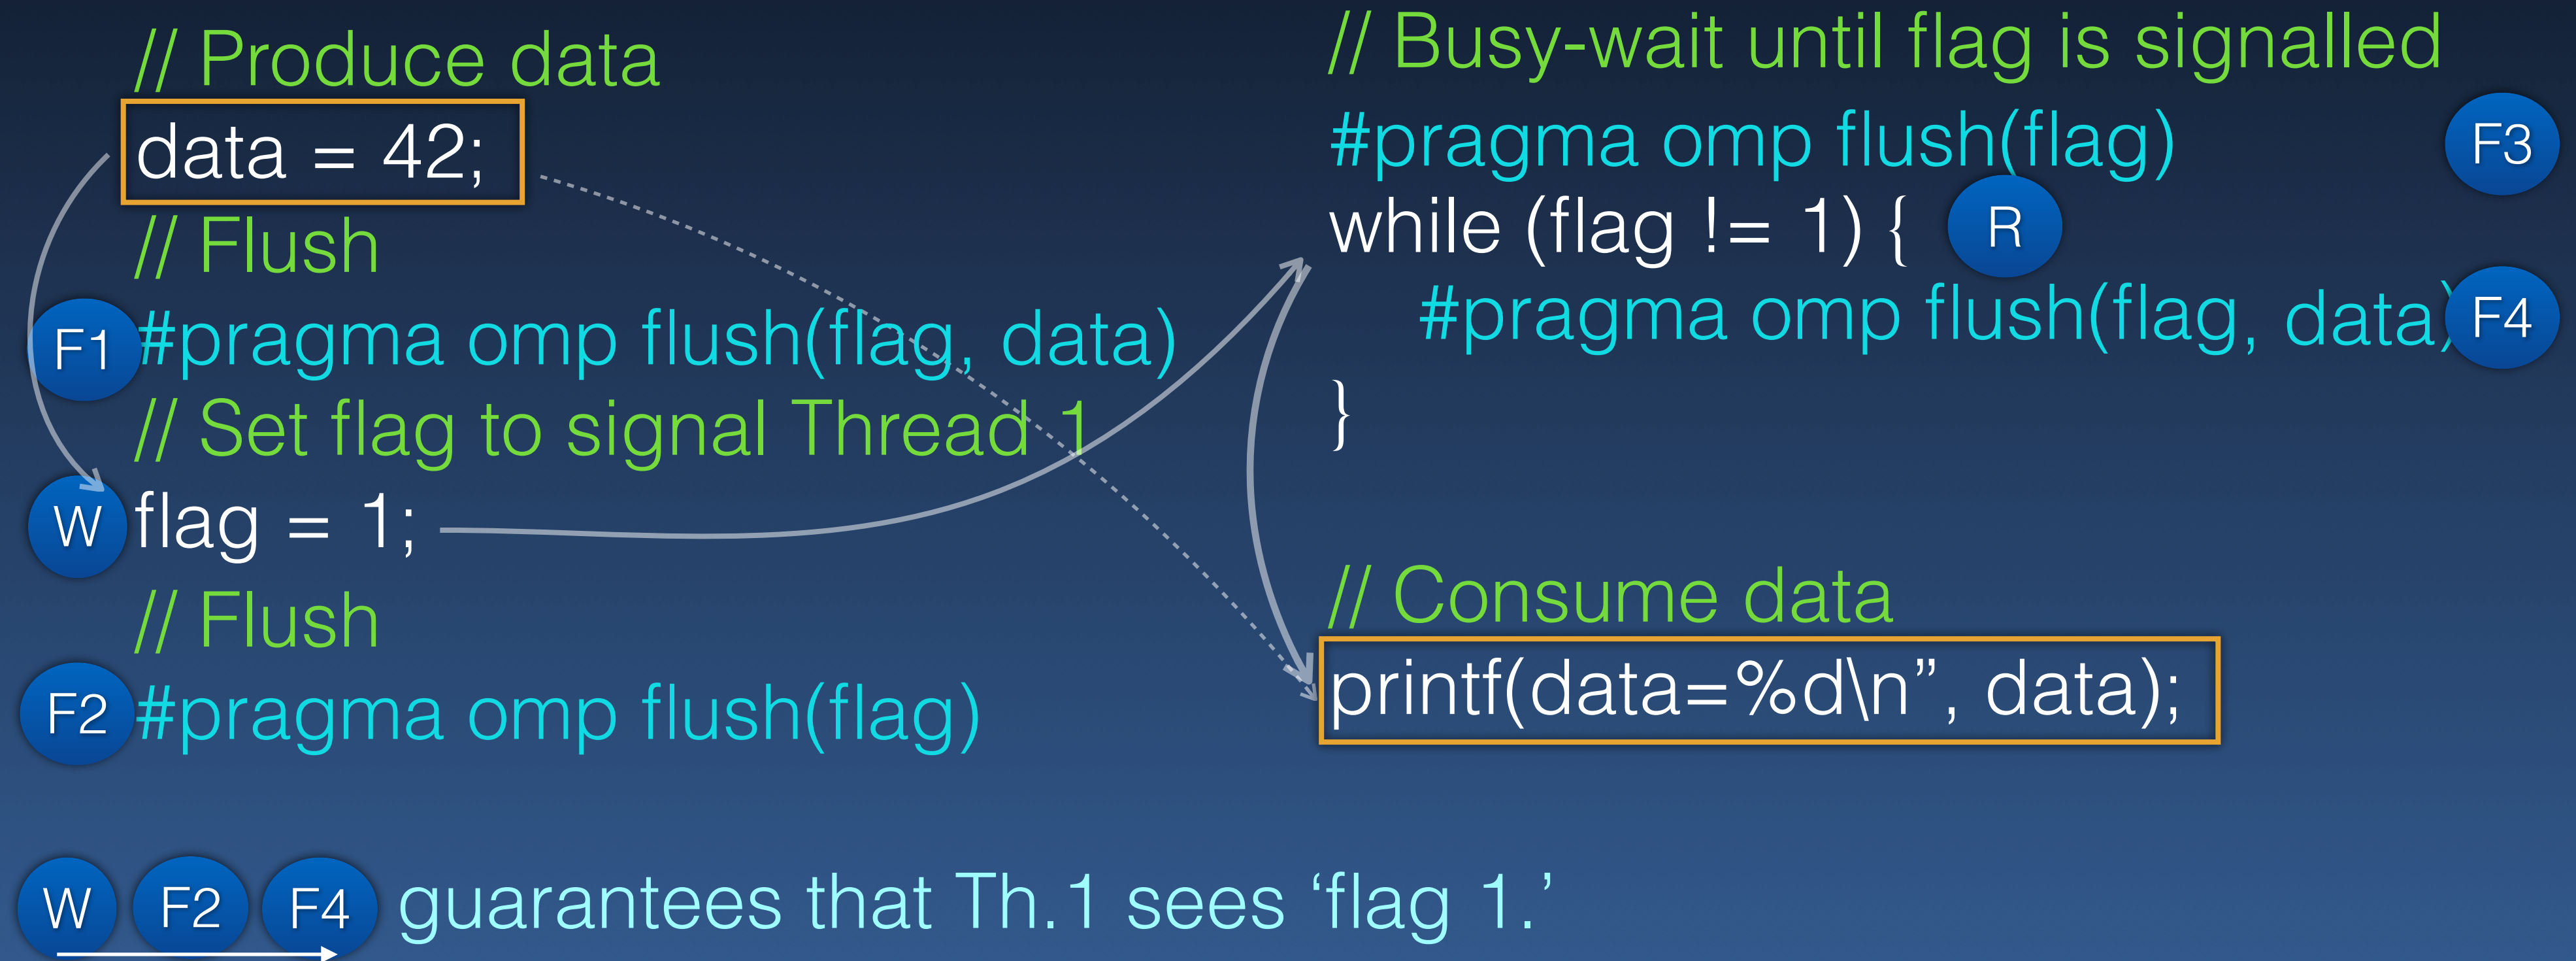


Producer-Consumer

int data, flag = 0;

Thread 0

Thread 1



Producer-Consumer

int data, flag = 0;

Thread 0

Thread 1

// Produce data
data = 42;
// Flush
F1 #pragma omp flush(flag, data)
// Set flag to signal Thread 1
W flag = 1;
// Flush
F2 #pragma omp flush(flag)

// Busy-wait until flag is signalled
#pragma omp flush(flag) F3
while (flag != 1) { R
#pragma omp flush(flag, data) F4
}

// Consume data
printf(data="%d\n", data);

W F2 F4 guarantees that Th.1 sees 'flag 1.' ('flag 1' \nRightarrow F2 \rightarrow F4)

Assume F2 must eventually finish.

Producer-Consumer

int data, flag = 0;

Thread 0

```
// Produce data
data = 42;
// Flush
F1 #pragma omp flush(flag, data)
// Set flag to signal Thread 1
W flag = 1;
// Flush
F2 #pragma omp flush(flag)
```

Thread 1

```
// Busy-wait until flag is signalled
#pragma omp flush(flag) F3
while (flag != 1) { R
    #pragma omp flush(flag, data) F4
}
// Consume data
printf(data="%d\n", data);
```

W F2 F4 guarantees that Th.1 sees 'flag 1.' ('flag 1' \nRightarrow F2 \rightarrow F4)

Assume F2 must eventually finish.

Producer-Consumer

int data, flag = 0;

Thread 0

```
// Produce data
data = 42;
// Flush
F1 #pragma omp flush(flag, data)
// Set flag to signal Thread 1
W flag = 1;
// Flush
F2 #pragma omp flush(flag)
```

Thread 1

```
// Busy-wait until flag is signalled
#pragma omp flush(flag) F3
while (flag != 1) { R
    #pragma omp flush(flag, data) F4
}

// Consume data
printf(data="%d\n", data);
```

W F2 F4 guarantees that Th.1 sees 'flag 1.'

'flag 1' in Th.1 \Rightarrow W has started and hence F1 has happened \Rightarrow F1 R

Producer-Consumer

int data, flag = 0;

Thread 0

```
// Produce data
data = 42;
// Flush
F1 #pragma omp flush(flag, data)
// Set flag to signal Thread 1
W flag = 1;
// Flush
F2 #pragma omp flush(flag)
```

Thread 1

```
// Busy-wait until flag is signalled
#pragma omp flush(flag) F3
while (flag != 1) { R
    #pragma omp flush(flag, data) F4
}
#pragma omp flush(flag, data) F5
// Consume data
printf(data="%d\n", data);
```

W F2 F4 guarantees that Th.1 sees 'flag 1.'

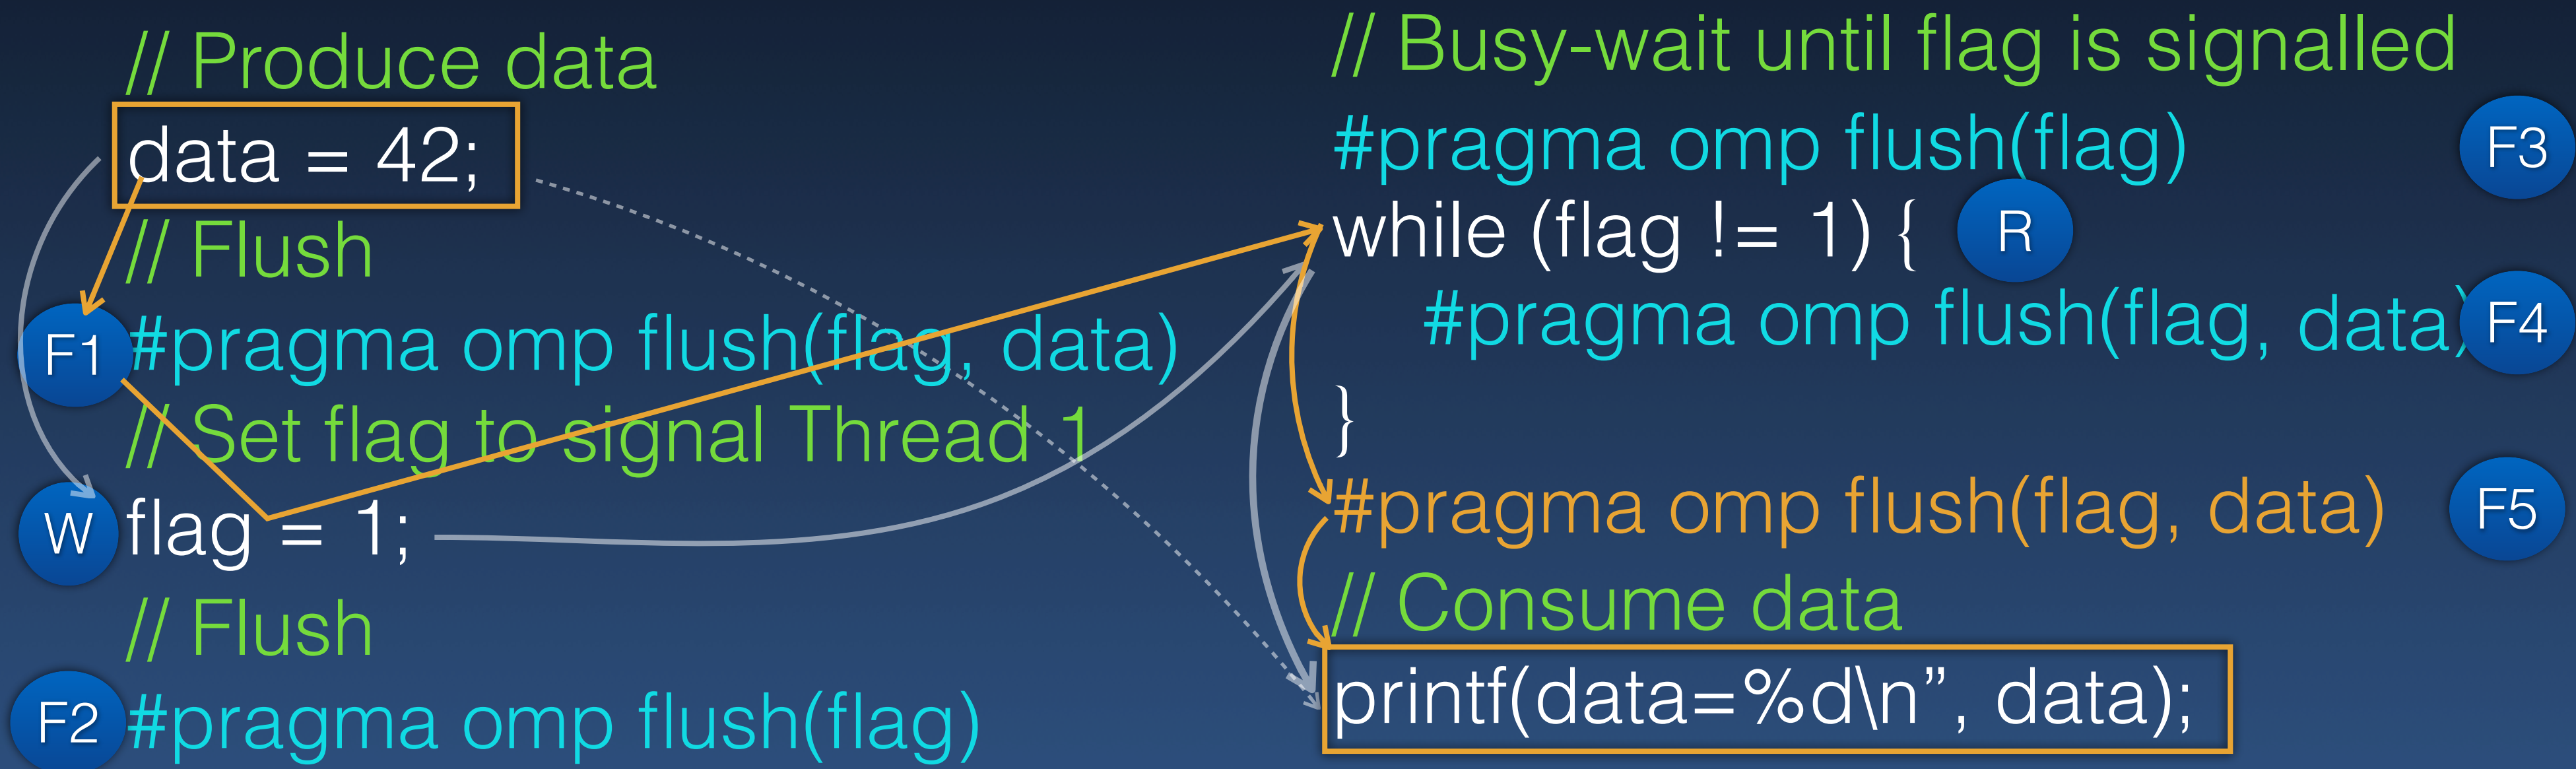
'flag 1' in Th.1 \Rightarrow W has started and hence F1 has happened \Rightarrow F1 R F5

Producer-Consumer

int data, flag = 0;

Thread 0

Thread 1



`W` `F2` `F4` guarantees that Th.1 sees 'flag 1.'

'flag 1' in Th.1 \Rightarrow `W` has started and hence `F1` has happened \Rightarrow `F1` `R` `F5`

Consistency Summary

Model	Description
Strict	Global time based atomic ordering of <i>all</i> shared accesses
Sequential	<i>All</i> threads see all shared accesses in the same order consistent with program order -- no centralized ordering
Causal	All threads see causally-related shared accesses in the same order
Processor	All threads see writes from each other in the order they were made. Writes to a variable must be seen in the same order by all threads
Weak	Special synchronization based reordering -- shared data consistent only after synchronization

- Causal, Processor, FIFO consistency
- Weak consistency
- OpenMP Flush